

Reproducibility Budgets for Open-Weight LLM Evaluation: A Generalizability Theory Study of 7–9B Instruction-Tuned Models

Anonymous ACL submission

Abstract

Large Language Model (LLM) evaluation results vary across reproductions, but variance sources interact, and their relative importance shifts with the benchmark format. We apply Generalizability theory (G-theory) in a fully crossed design with up to seven facets across open-weight instruction-tuned models and five benchmarks. Under binary correctness scoring, a recommended 3-prompt $\times$ 3-seed design achieves $G \geq 0.80$ with 178 items—a $3.1\times$ item reduction over single-condition designs at the cost of $2.9\times$ more total evaluations—a tradeoff that favors multi-facet designs when item construction is costly. Variance decomposition reveals that models differ primarily in which items they find difficult rather than in overall ability, and that the variance structure shifts qualitatively between multiple-choice and free-form benchmarks. Preliminary validation across model scales and architecture families confirms structural stability. We release an open-source G-theory pipeline and D-study budget calculator for evaluation design optimization.

1 Introduction

In our analysis of five benchmarks, item selection drives 34–38% of evaluation variance on multiple-choice tasks while prompt wording is negligible ($<3\%$). On the free-form reasoning task GSM8K, the structure shifts: model $\times$ prompt interaction alone accounts for approximately 9%, while item selection drops to 17%. These numbers, drawn from a factorial analysis of approximately 2.3 million evaluation records, reveal that the variance structure of LLM evaluation is not a fixed property of the pipeline but shifts qualitatively with the response format. Current evaluation practice treats variance sources as independent nuisance factors and controls them one at a time: one study perturbs prompts, another varies precision, a third changes

item subsets. This piecemeal approach cannot capture interactions between sources, nor can it predict how the relative importance of each source changes across evaluation settings. Leaderboard rankings built on single-condition evaluations are therefore sensitive to methodological choices that remain unreported and uncontrolled (Rodriguez et al., 2021; Dehghani et al., 2021). No study has simultaneously crossed infrastructure and methodology facets in a unified factorial design across benchmarks to decompose these sources, quantify their interactions, and compare the resulting variance structure.

Existing studies examine individual variance sources in isolation. Sclar et al. (2024) and BrittleBench (Romanou et al., 2026) documented 50–76% accuracy swings from prompt formatting; Yuan et al. (2025) showed up to 9% variation from numerical precision; Pezeshkpour and Hruschka (2024), Dodge et al. (2020), and Henderson et al. (2018) documented ordering and hyperparameter sensitivity—each holding other factors fixed. Messing (2026) applied G-theory to LLM evaluation across five methodology facets on a single benchmark (judge identity: 43.8%), but excluded infrastructure facets and cross-benchmark comparison. Madaan et al. (2024) proposed per-factor metrics without crossing factors. IRT approaches (Maia Polo et al., 2024; Lalor et al., 2019; Zhou et al., 2026) optimize item selection but ignore prompt, seed, and ordering facets. Bouthillier et al. (2021) established ANOVA-based variance decomposition for ML reproducibility. Three gaps remain: no study has (a) crossed infrastructure and methodology facets in a unified factorial design, (b) compared variance structures across response formats, or (c) translated variance components into design-dependent reproducibility budgets.

Through a factorial G-theory study (Shavelson and Webb, 1991; Brennan, 2001) crossing up to seven facets across eight open-weight 7–9B param-




Figure 1: Overview of the G-theory evaluation pipeline. A factorial design crossing seven facets across eight models and five benchmarks produces 2.3M binary records. G-study variance decomposition via Henderson Method I reveals that model×item interaction accounts for up to 36% of variance. D-study projections translate variance components into design-dependent item budgets: a 3-prompt × 3-seed design achieves $G \geq 0.80$ with 178 items (3.1× reduction). Two contributions emerge: reproducibility budgets (C1) and format-associated variance patterns (C2).




Figure 2: D-study projections for MMLU (single-model baseline). A recommended 3-prompt × 3-seed design reduces the budget from 552 to 178 items for $G \geq 0.80$; the full multi-facet design reduces further to 55 items.

eter models and five benchmarks (MMLU, ARC-Challenge, HellaSwag, GSM8K, MATH), totaling approximately 2.3 million records, we find that evaluation variance has a stable but format-dependent internal structure. D-study projections show that a recommended 3-prompt × 3-seed design achieves $G \geq 0.80$ with 178 items (1,602 evaluations per model), a 3.1× reduction from 552 single-condition items; the full 1,296-condition design further reduces to 55 items as a theoretical ceiling (Figure 2). Model×item interaction (up to 36%) indicates that models differ primarily in *which items* they find difficult (Pezeshkpour and Hruschka, 2024; Maia Polo et al., 2024).

Our contributions:

1. Design-dependent reproducibility budgets.

A recommended 3-prompt × 3-seed design achieves $G \geq 0.80$ with 178 items (1,602

evaluations per model), a 3.1× reduction over single-condition designs (552 items); the full 1,296-condition design reduces to 55 items (71,280 evaluations). Under binary scoring, BF16 and FP16 require no separate control (<0.1% interaction). IRT item selection and G-theory protocol optimization compose: IRT-selected 116 items yield $G = 0.457$ single-condition but $G = 0.895$ multi-facet (Table 4).

2. Format-associated variance patterns (exploratory).

MC benchmarks concentrate 68–72% of variance in item-related components; FF benchmarks disperse to 31–41% (item-level $d = 1.73$; benchmark-level confirmation limited by sample size). This hypothesis-generating observation warrants confirmation with larger benchmark samples.

We release an open-source G-theory pipeline and D-study budget calculator for evaluation design optimization.

2 Related Work

Variance Decomposition in ML/NLP Evaluation. Factorial decomposition of evaluation variance has roots in the ML reproducibility literature. Bouthillier et al. (2021) introduced an ANOVA-based template that decomposes benchmark variance across implementation details, hyperparameters, and data splits, establishing that design choices account for a substantial share of reproducibility failures. Subsequent studies applied similar decomposition to specific evaluation

facets (Dodge et al., 2020; Henderson et al., 2018; Magnusson et al., 2023). Messing (2026) brought Generalizability theory into LLM evaluation, decomposing variance across five methodology facets on a single benchmark (judge identity: 43.8%; see Positioning below for detailed comparison). Madaan et al. (2024) proposed per-factor variance metrics for LLM evaluation benchmarks but did not cross factors simultaneously. Holistic evaluation frameworks (Liang et al., 2023; Kiela et al., 2021) standardize multi-scenario comparison but do not decompose variance sources. Blackwell et al. (2024) quantified uncertainty in benchmark scores, and Qian et al. (2026) conducted systematic meta-evaluation of LLM benchmarks; Burnell et al. (2023) called for rethinking evaluation reporting standards; Saxon et al. (2024) advocated model metrology as a framework for rigorous benchmark design. Recently, G-theory has been applied to LLM-based writing assessment (Song et al., 2025) and annotation reliability (Camuffo et al., 2026), but these studies address scoring consistency rather than benchmark evaluation variance. Kunievsky and Evans (2025) proposed an alternative decomposition that separates intent uncertainty, articulation variance, and model uncertainty, but this framework does not cross infrastructure and methodology facets in a factorial design. Our work extends these approaches by crossing seven facets across five benchmarks, enabling direct comparison of variance structures across evaluation formats.

Infrastructure and Prompt Variance. A parallel line of work examines individual variance sources in isolation. Yuan et al. (2025) documented up to 9% accuracy variation from numerical precision; Messina and Scotta (2026) and He and Thinking Machines Lab (2025) formalized and mitigated CUDA nondeterminism from floating-point non-associativity. On the methodology side, Mizrahi et al. (2024) advocated multi-prompt evaluation based on large-scale sensitivity analysis; Sclar et al. (2024) and BrittleBench (Romanou et al., 2026) found 50–76% accuracy swings from prompt formatting; Hua et al. (2025) showed that much of this sensitivity is an artifact of log-likelihood scoring (our generation-based protocol controls for this). Pezeshkpour and Hruschka (2024) and Alzahrani et al. (2024) documented ordering and benchmark sensitivity; Bui et al. (2025) and Biderman et al. (2024) documented seed and training sensitivity respectively. Each isolates one factor, precluding

interaction estimation. Standardized evaluation frameworks such as lm-evaluation-harness (Gao et al., 2024) reduce implementation variance but do not control prompt, seed, or ordering facets.

IRT and Efficient Benchmarking. Item Response Theory (IRT) offers a complementary perspective on evaluation design by modeling item difficulty and discrimination parameters. Lalor et al. (2019) introduced IRT-based item analysis for NLP benchmarks, and Rodriguez et al. (2021) showed that IRT reveals hidden properties of evaluation items that raw accuracy scores obscure. Building on this foundation, Maia Polo et al. (2024) developed tinyBenchmarks, using IRT-based subsampling to reduce evaluation cost while preserving ranking accuracy, and Zhou et al. (2026) recommended a 116-item subset via PSN-IRT for efficient MMLU evaluation. Perlitz et al. (2023) and Kipnis et al. (2025) explored related strategies for benchmark compression; Li et al. (2025) extended IRT to adaptive testing. Complementary benchmark revision efforts address item quality directly: MMLU-Redux (Gema et al., 2025) identified annotation errors in MMLU, while MMLU-Pro (Wang et al., 2024) extended it with harder items and expanded option sets. These approaches optimize item selection or quality but do not model variance from prompt, seed, or ordering facets. IRT-selected items evaluated under a single methodological condition can therefore systematically underestimate the number of items required for reliable measurement across diverse evaluation settings.

Psychometrics Foundations. G-theory (Shavelson and Webb, 1991; Brennan, 2001; Cardinet et al., 2010) provides the canonical framework for multifacet reliability analysis; Webb et al. (2006) detail the random versus fixed facet distinctions central to our model specification. Extensions to binary data include Mushquash and O’Connor (2006) and logistic-link generalized G-theory (Jiang and Skorupski, 2018; Tuerlinckx et al., 2006). We operationalize differential item functioning—established as a formal framework by Holland and Wainer (1993) and operationalized via statistical detection procedures by Clauser and Mazor (1998)—as model×item interaction. Recent psychometric perspectives on LLM evaluation address construct validity (Kearns, 2026; Bean et al., 2025; Freiesleben, 2026) and measurement properties (Ye et al., 2025); Zhu et al. (2026) propose a benchmark health index for systematic meta-evaluation.

Positioning. Closest to our work, [Messing \(2026\)](#) applied G-theory to LLM evaluation through their TEE framework, decomposing five methodology facets on a single benchmark (judge identity: 43.8%). We extend this in three directions: (i) we fix the judge to the model itself, replacing judge identity with infrastructure facets and additional methodology facets (seed, ordering); (ii) we cross facets across five benchmarks, enabling cross-format comparison; and (iii) we translate variance components into reproducibility budgets via D-study projections. The two studies are complementary: TEE characterizes judge selection, while we characterize item-and-protocol selection. Relative to IRT approaches ([Maia Polo et al., 2024](#); [Zhou et al., 2026](#)), G-theory optimizes an orthogonal dimension: conditions needed for target reliability rather than item discrimination.

3 Method and Experimental Design

3.1 Generalizability Theory

Generalizability theory (G-theory) treats each measurement as one observation from a universe of admissible conditions defined by the crossed facets of an evaluation design ([Shavelson and Webb, 1991](#); [Brennan, 2001](#)). The *object of measurement* is the entity whose consistency is of interest. In our setting, test items (p) serve as the object of measurement, and the goal is to determine how reliably item-level scores generalize across the methodological conditions under which they are collected.

A *G-study* decomposes total observed variance into components σ_x^2 for each facet x and their interactions. Each component quantifies how much variability that source contributes to the overall measurement. Large interaction components (e.g., σ_{pM}^2 for item-by-model) indicate that facets do not affect items uniformly but instead create differential difficulty profiles. We estimate variance components via Henderson Method I, an unbiased ANOVA-based estimator for balanced designs ([Brennan, 2001](#), Ch. 3), with 1,000-resample bootstrap confidence intervals ([Efron and Tibshirani, 1993](#)) (item-level resampling with replacement, preserving the fully-crossed structure across all other facets) ([Li, 2023](#)).

A *D-study* projects the generalizability coefficient under hypothetical designs with different numbers of conditions per facet:

$$G_{\text{item}} = \frac{\sigma_p^2}{\sigma_p^2 + \sigma_{\text{rel}}^2} \quad (1)$$

where $\sigma_{\text{rel}}^2 = \sum_i \sigma_{p \times f_i}^2 / n_{f_i} + \sigma_{\text{residual}}^2 / \prod_i n_{f_i}$ aggregates all item-by-facet interaction variances, each divided by the number of conditions for the corresponding facets ([Brennan, 2001](#), Ch. 4). A D-study projects reliability for a test averaging over n_p items via

$$G(n_p) = \frac{n_p \sigma_p^2}{n_p \sigma_p^2 + \sigma_{\text{rel}}^2} \quad (2)$$

enabling independent verification of all budget projections in Table 4. This prospective capability distinguishes G-theory from classical ANOVA: rather than only attributing variance post hoc, a D-study enables optimization of evaluation designs *before* data collection ([Shavelson et al., 1989](#)). We define *item-related variance* as the sum of the item main effect (σ_p^2) and the model $\times$ item interaction (σ_{pM}^2). The latter captures differential item functioning (DIF): items whose difficulty ranking changes across models ([Brennan, 2001](#), Ch. 4). Because DIF is operationally an item-level property—it reflects model-specific item difficulty profiles rather than a global model ability shift—aggregating $\sigma_p^2 + \sigma_{pM}^2$ provides a unified measure of item-driven variance. We report unbundled components throughout to ensure transparency.¹

3.2 Experimental Design: MMLU

We design a fully crossed factorial experiment (exp-001) that simultaneously varies seven facets of LLM evaluation on MMLU. Table 1 summarizes the complete design matrix. Figure 1 illustrates the complete pipeline from factorial design through variance decomposition to budget projection.

Facets. Items (p): 200 MMLU questions selected via purposive stratified sampling across 10 of MMLU’s 57 subjects (20 per subject, deterministic selection with arbitrarily chosen seed 42). The 10 subjects span four disciplinary domains—STEM: abstract algebra, college physics, high school mathematics, computer security; humanities: world religions, philosophy; social sciences: US foreign policy, international law; professional: professional medicine, clinical knowledge—ensuring that variance estimates are not dominated by a single knowledge domain. Resampling stability (CV = 3.1%) was evaluated within these 10 subjects; cross-subject generalization remains untested.

¹Notation: p = item, M = model, π = precision, δ = prompt, s = seed, τ = temperature, o = ordering, σ_x^2 = variance component for source x .

Table 1: Fully crossed G-theory design. Exp-001 evaluates MMLU across all facet levels; exp-002 extends to five benchmarks with reduced crossing. [†]BF16 balanced subset (2 temperatures: 0.0, 0.7) for cross-model analysis; the full single-model design uses all 3 levels.

Facet	Levels	exp-001	exp-002
Item (p)	Stratified sample	200	200 $\times$ 5
Model (M)	8 architecture families	8	8
Precision (π)	BF16 / FP16 / FP32	3	1
Prompt (δ)	Paraphrases	6	6
Seed (s)	Fixed set	6	6
Temperature (τ)	0.0 / 0.3 / 0.7	3	2
Ordering (o)	Demo permutations	4	4
Total records		460,800 [†]	2,304,000

Models (M): eight open-weight 7–9B parameter instruction-tuned models spanning eight architecture families—DeepSeek-7B, Gemma-2-9B, InternLM2.5-7B, Llama-3.1-8B-Instruct, Mistral-7B-v0.3-Instruct, OLMo-2-7B, Qwen3-8B, and Yi-1.5-9B (full identifiers in Appendix H)—treated as a quasi-random facet (§3.6). Numerical precision (π): BF16, FP16, and FP32, all evaluated with `enforce_eager=True` for fair comparison across precision levels. Prompt template (δ): six surface perturbations of the evaluation instruction (Sclar et al., 2024), varying option layout (lettered vs. numbered), instruction wording (direct vs. contextual), and answer prompt cue for MC benchmarks; for FF benchmarks, templates vary chain-of-thought elicitation strategy (full template text in Appendix H). Random seed (s): six fixed values (42, 123, 456, 789, 1024, 2048). Sampling temperature (τ): three levels (0.0, 0.3, 0.7) spanning deterministic to moderately stochastic decoding. Answer ordering (o): four permutations of few-shot demonstration examples (for MMLU; structurally inert for other benchmarks, see §3.3) (Pezeshkpour and Hruschka, 2024).

Sample sizes. The cross-model BF16 balanced subset yields 460,800 records (8 models $\times$ 200 items $\times$ 6 prompts $\times$ 6 seeds $\times$ 2 temperatures $\times$ 4 orderings). A single-model analysis (Llama) spans all three precisions, yielding 259,200 binary observations across 1,296 fully balanced conditions. The BF16 subset serves as the primary cross-model sample because BF16 is the standard inference precision for these architectures.

3.3 Cross-Benchmark Extension

To test whether the variance structure observed on MMLU generalizes across evaluation formats, exp-002 extends the design to five benchmarks: MMLU,

ARC-Challenge, and HellaSwag (multiple-choice), and GSM8K and MATH (free-form generation). The design crosses 8 models $\times$ BF16 $\times$ 2 temperatures (0.0, 0.7) $\times$ 6 prompts $\times$ 6 seeds $\times$ 4 orderings $\times$ 200 items per benchmark, totaling 2,304,000 records. This reduced crossing fixes precision at BF16 and uses two temperature levels to concentrate statistical power on format-associated differences. The ordering facet operates only on MMLU, where it selects among four permutations of the few-shot demonstration examples. For all other benchmarks, the prompt construction does not use the ordering parameter; ordering is structurally inert ($\sigma_{\text{ordering}}^2 = 0$ by construction) and is retained solely for design balance. The two temperature levels bracket the most common evaluation configurations: deterministic ($\tau = 0.0$) and a moderate stochastic setting ($\tau = 0.7$) representative of practical deployment. A subset analysis confirms that this reduced crossing preserves the variance structure of the full design ($\Delta G = 0.0008$; Appendix A).

3.4 Evaluation Protocol

All evaluations use generation-based scoring (free-text response with answer extraction) rather than log-likelihood ranking, following Hua et al. (2025). Binary correctness (1/0) is the primary metric; text-exact-match is secondary. Inference uses vLLM 0.9.2 without batch-invariant kernels, preserving CUDA-level nondeterminism as a measurable variance source (Messina and Scotta, 2026; He and Thinking Machines Lab, 2025). Few-shot protocols: 5-shot MMLU/ARC, 4-shot MATH, 8-shot GSM8K, 10-shot HellaSwag.

3.5 Binary Data Assumptions

Standard G-theory assumes continuous, approximately normal outcome data (Shavelson and Webb, 1991). Our binary correctness scores (0/1) violate this assumption. We establish robustness through simulation and estimator properties, with two additional analyses (GLMM complement, accuracy-range stability) in Appendix A.

Henderson Method I produces unbiased estimates for balanced designs regardless of outcome distribution (Brennan, 2001, Ch. 12), and REML estimates differ by <0.5pp on all components. A probit simulation, GLMM complement ($\rho = 1.0$), and accuracy-range analysis all confirm robustness (Appendix A).

All findings in this paper are conditioned on binary correctness scoring. This measurement resolution collapses continuous generation quality into a pass/fail outcome, which may amplify item-related variance by reducing within-item response variability and mask finer-grained variance sources (e.g., partial credit, reasoning quality) that would emerge under richer scoring rubrics. The text-exact-match comparison in §4.1 ($G = 0.007$ vs. 0.936) illustrates how the scoring metric can fundamentally reshape variance structure. A partial-credit analysis on GSM8K (8 models, step-level scoring) confirms this conditionality: G_{item} drops from 0.86 to 0.77 , with all eight models showing $G_{\text{partial}} < G_{\text{binary}}$ (mean $\Delta G = -0.135$; Appendix A), providing evidence that D-study budgets have limited transferability to finer-grained scoring rubrics.

3.6 Model Facet Specification

Our target universe comprises open-weight, instruction-tuned models in the 7–9B parameter range (eight architecture families); results may not generalize to API-based, base, or >9B models. Our eight purposively selected models are treated as a quasi-random facet (Meredith, 1993; Putnick and Bornstein, 2016). We report G under random-model assumptions, the more conservative specification: across five benchmarks, treating models as fixed yields G_{fixed} of 0.951 – 0.990 (mean 0.975) versus G_{random} of 0.863 – 0.896 (mean 0.882), a consistent increase of 0.077 – 0.113 (mean 0.093), because model $\times$ item interaction (30–36%) enters the relative error variance under random assumptions. Consequently, our D-study budgets are upper bounds—practitioners will not under-sample regardless of facet classification. The variance decomposition itself (Tables 2–3) is invariant to this choice. Leave-one-model-out analysis preserves component rankings across all five benchmarks, and Cohen’s κ pairwise agreement (mean $\kappa = 0.33$ – 0.46) confirms genuine differential item functioning (Appendix A). Four-model subset consistency checks match the observed eight-model G within mean $|\Delta G| = 0.035$ across all five benchmarks (§4.5).

For reference, G_{model} (treating models as the object of measurement) under the recommended 3-prompt $\times$ 3-seed design with 200 items averages 0.79 across five benchmarks (range: 0.52 on ARC to 0.93 on HellaSwag). Model ranking reliability depends primarily on model $\times$ prompt interaction—ARC (8.0%) and GSM8K (8.5%) show lower G_{model} because prompt sensitivity reduces model

discriminability. G_{item} remains the primary metric because item reliability is the direct optimization target for evaluation design; model reliability is a downstream consequence that improves with item count.

4 Experiments

All experiments use open-weight 7–9B parameter models evaluated with vLLM 0.9.2 on a single NVIDIA RTX PRO 6000 GPU (96GB) with binary correctness scoring.

4.1 Single-Model Baseline

We first establish the variance structure within a single model (Llama-3.1-8B-Instruct) on MMLU to isolate methodological facet contributions before introducing cross-model variation.

Design. The G-study crosses six facets: 200 items $\times$ 3 numerical precisions (BF16, FP16, FP32) $\times$ 6 prompt templates $\times$ 6 random seeds $\times$ 3 temperature levels ($0.0, 0.3, 0.7$) $\times$ 4 answer orderings, yielding 259,200 binary observations. Variance components are estimated via Henderson Method I with bootstrap confidence intervals (1,000 resamples).

Results. Under binary scoring, item difficulty (item_id) accounts for 58.4% of total variance, with the residual term contributing 23.1% . All main facet effects are negligible ($<1\%$ each): precision 0.9% , prompt template, temperature, seed, and ordering each $<0.5\%$. The interesting structure appears in item-conditional interactions: prompt $\times$ item is the largest at 5.5% , followed by seed $\times$ item (4.7%), precision $\times$ item (3.3%), temperature $\times$ item (2.9%), and ordering $\times$ item (0.8%). This pattern indicates that methodological choices do not shift performance uniformly across items; instead, specific items are differentially sensitive to specific facets—prompt template variation produces the largest item-conditional effect.

The resulting generalizability coefficient is $G_{\text{item}} = 0.936$, indicating that 200 items under multi-facet conditions produce highly reliable rankings. A D-study projection shows $G \geq 0.80$ is achievable with as few as 55 items in this multi-facet design.

For text-exact-match scoring (requiring verbatim string matches), the variance structure shifts dramatically: temperature dominates at 79.7% of variance, and G_{item} collapses to 0.007 . This metric-

Table 2: Variance decomposition for the cross-model MMLU design (8 models, BF16, 460,800 records). Components estimated via Henderson Method I. $G_{\text{item}} = 0.896$.

Component	Var. %
item_id	37.9
<u>model×item</u>	<u>30.2</u>
residual	25.0
model	2.8
model×prompt	1.6
prompt×item	0.9
temp.×item	0.6
seed×item	0.5
ordering×item	0.3
all others (each)	<0.1

conditional divergence motivates our exclusive use of binary correctness for the remaining analyses.

4.2 Cross-Model Variance

Extending to eight models (BF16 subset, 460,800 records), Table 2 shows that item difficulty drops from 58.4% to 37.9%, while model×item interaction emerges at 30.2% (up to 36% across benchmarks, §4.3). The model main effect is 3–5% (MC) and 9–11% (FF)², confirming that models differ primarily in *which items* they find difficult rather than in overall ability level (contamination implications in §6).

4.3 Exploratory: Format-Associated Variance Patterns

We extend to five benchmarks (three MC: MMLU, ARC-Challenge, HellaSwag; two FF: GSM8K, MATH) using exp-002 (8 models × BF16 × 2 temperatures × 6 prompts × 6 seeds × 4 orderings × 200 items). Table 3 reveals a format-associated pattern: MC benchmarks concentrate item-related variance (item + model×item DIF) at 68–72%, while FF benchmarks disperse it to 31–41% ($n = 5$; permutation $p = 0.10$; item-level $d = 1.73$, benchmark-level CI spans zero); we report this as a descriptive observation. This does not extend to item budgets: HellaSwag (MC, 118 items for $G \geq 0.80$) exceeds ARC (98) and MATH (101), showing format alone does not predict evaluation cost (Figure 3). The factorial design is not fully symmetric: ordering applies only to MMLU, and prompt templates serve format-specific functions, potentially confounding format comparisons. Per-benchmark details—including HellaSwag and

Table 3: Format-associated variance structure across five benchmarks (exp-002, 8 models). Variance components (%) per benchmark; MC = multiple-choice, FF = free-form. Bold = dominant non-residual component. [†]For MC, reports prompt×item interaction (prompt main effect <1%); for FF, reports prompt main effect (item-conditional interactions absorbed into residual).

	MC			FF	
	MMLU	ARC	Hella.	GSM	MATH
item	37.9	37.9	34.3	16.9	24.0
mod.×item	30.2	33.7	36.0	14.6	17.3
model	2.8	3.1	5.4	10.7	8.5
prompt [†]	0.9	1.4	2.1	5.4	0.8
mod.×prmt	1.6	8.0	0.2	8.5	1.7
residual	25.0	15.4	21.4	39.9	43.9
all others	1.6	0.5	0.6	4.0	3.8




Figure 3: Variance decomposition (Henderson Method I) across five benchmarks and eight models. Multiple-choice benchmarks share an item-dominated structure; HellaSwag is the exception where model×item (36.0%) exceeds item difficulty (34.3%). Free-form benchmarks (GSM8K, MATH) distribute variance across item, model, and interaction components with no single dominant source.

ARC exceptions, FF residual analysis (Table 11), difficulty dispersion as a variance predictor (Figure 5; Table 10)—are in Appendix D. All format-associated findings are conditional on binary scoring (§3.5).

4.4 Precision Separability

Under binary scoring, BF16 and FP16 are functionally interchangeable (<0.1% interaction (Yuan et al., 2025)); under text-exact-match, a precision×temperature interaction emerges at 1.3%, consistent with §4.1. Precision need not be controlled for binary evaluation at BF16/FP16 (Shanmugavelu et al., 2024).

4.5 Design-Dependent Reproducibility Budgets

How many items does a given evaluation design need to achieve a target reliability? We translate variance components into item budgets via D-study projections.

²Ranges rounded to nearest integer; exact values are 2.8–5.4% (MC) and 8.5–10.7% (FF) per Table 3.

Table 4: D-study reproducibility budgets under different evaluation designs (single-model, MMLU). Total evals = items $\times$ conditions.

Design	Cond.	$n_{G \geq .80}$	Total	G_{200}	G_{116}
Multi-facet	1,296	55	71,280	.936	—
Single-cond.	1	552	552	.592	.457
IRT+multi	1,296	—	150,336	—	.895 [‡]
IRT baseline	—	116	116	—	.457 [†]

[†] PSN-IRT 116 items (Zhou et al., 2026) under single-condition: $G = 0.457$.

[‡] Same 116 items under multi-facet: $G = 0.895$.

Table 4 compares evaluation designs at the $G \geq 0.80$ threshold using single-model (Llama-3.1-8B-Instruct) MMLU variance components. The multi-facet design (1,296 conditions per item) achieves $G = 0.936$ at 200 items and requires only 55 items for $G \geq 0.80$. A single-condition design reaches only $G = 0.592$ at 200 items and requires 552 items—a 10.0 $\times$ budget ratio stable across subset sizes (CV = 1.9%).

Zhou et al. (2026)’s IRT-selected 116 items yield $G = 0.457$ under single-condition but $G = 0.895$ under multi-facet conditions, demonstrating that IRT item selection and G-theory protocol optimization are orthogonal and complementary. An internal consistency check using four-model subsets confirms projection stability ($|\Delta G| \leq 0.029$ for MC, ≤ 0.065 for FF; Table 8); the corresponding $\sim 23\%$ budget increase is consistent across all five benchmarks (Figure 4).

The design choice depends on the evaluation bottleneck: the full 1,296-condition design minimizes items (55, 71,280 evaluations/model); the recommended 3×3 design balances cost and reliability (178 items, 1,602 evaluations); single-condition requires 552 items. Cross-model projections show budget ratios of 11–22 $\times$ across benchmarks (bootstrap CIs all exceed 10 $\times$; Appendix B), with multi-facet budgets of 93–128 items (Table 7).

Projection assumptions. Cross-model budgets increase to 93 items because model $\times$ item interaction enters relative error variance. G_{item} is stable across benchmarks (0.863–0.896, CV = 1.4%; LOMO in Appendix A). All projections assume the 2024–2025 model generation; the 72B extension (§4.6) provides preliminary support.

4.6 Preliminary Scale Validation

Within-family analysis (Qwen2.5 3B/7B/72B on MMLU) and a cross-family check (Llama-3.1-70B-Instruct) confirm qualitatively invariant variance structure: item difficulty dominates at all scales

(73–83%, all single-model $G > 0.95$), with consistent facet ranking across a 24 $\times$ parameter range and two architecture families (full results in Appendix A).

Robustness checks. Three MMLU resamples yield CV = 3.1%; treating models as fixed increases G by 0.077–0.113, confirming conservative budgets (§3.6; seven additional analyses in Appendix A).

5 Conclusion

A recommended 3-prompt $\times$ 3-seed design reduces item requirements from 552 to 178 for $G \geq 0.80$, confirmed across five benchmarks under binary scoring. The variance structure is format-dependent: MC benchmarks concentrate 68–72% of variance in item-related components, while FF benchmarks disperse to 31–41%, suggesting that reproducibility budgets should be calibrated per benchmark format. All budgets assume binary scoring; extending to other model scales and continuous metrics is a priority. We recommend that benchmark designers adopt multi-facet validation (at minimum 3 prompts $\times$ 3 seeds) as standard practice, and that leaderboard operators average over prompt and seed conditions to improve ranking stability. Code, data, and a D-study calculator will be released upon publication.

6 Limitations

Our design covers eight 7–9B instruction-tuned models from the 2024–2025 generation; preliminary scale validation with Qwen2.5-72B and Llama-3.1-70B confirms consistent variance structure across two families, but broader architectural and temporal coverage is needed. The benchmark sample ($n = 5$) limits statistical power for format-associated observations ($p = 0.10$); code generation and multi-turn formats remain untested. All D-study budgets assume binary correctness scoring; partial-credit analysis on GSM8K shows G drops from 0.86 to 0.77, suggesting limited transferability to finer-grained rubrics. Differential benchmark contamination across architectures could inflate model $\times$ item interaction, though consistent structures across three independently constructed MC benchmarks provide indirect evidence against contamination as the sole driver.

References

- Norah Alzahrani, Hisham Abdullah Alyahya, Yazeed Alnumay, Sultan Alrashed, Shaykhah Alsubaie, Yusef Almushaykeh, Faisal Mirza, Nouf Alotaibi, Nora Altwairesh, Areeb Alowisheq, M Saiful Bari, and Haidar Khan. 2024. [When benchmarks are targets: Revealing the sensitivity of large language model leaderboards](#). In *Proceedings of the Association for Computational Linguistics (ACL)*, pages 13787–13805.
- Andrew M. Bean, Ryan Othniel Kearns, Angelika Romanou, Franziska Sofia Hafner, Harry Mayne, Jan Batzner, Negar Foroutan, Chris Schmitz, Karolina Korgul, Hunar Batra, Oishi Deb, Emma Beharry, Cornelius Emde, Thomas Foster, Anna Gausen, María Grandury, Simeng Han, Valentin Hofmann, Lujain Ibrahim, and 23 others. 2025. [Measuring what matters: Construct validity in large language model benchmarks](#). In *Advances in Neural Information Processing Systems 38 (NeurIPS 2025), Datasets and Benchmarks Track*.
- Stella Biderman, Hailey Schoelkopf, Lintang Sutawika, Leo Gao, Jonathan Tow, Baber Abbasi, Alham Fikri Aji, Pawan Sasanka Ammanamanchi, Sidney Black, Jordan Clive, Anthony DiPofi, Julien Etxaniz, Benjamin Fattori, Jessica Zosa Forde, Charles Foster, Jeffrey Hsu, Mimansa Jaiswal, Won Ik Lee, Haonan Li, and 11 others. 2024. [Lessons from the trenches on reproducible evaluation of language models](#). *arXiv preprint arXiv:2405.14782*.
- Robert E. Blackwell, Jon Barry, and Anthony G. Cohn. 2024. [Towards reproducible llm evaluation: Quantifying uncertainty in llm benchmark scores](#). *arXiv preprint arXiv:2410.03492*.
- Xavier Bouthillier, Pierre Delaunay, Mirko Bronzi, Asya Trofimov, Brennan Nichyporuk, Justin Szeto, Nazanin Mohammadi Sepahvand, Edward Raff, KshitiJ Madan, Vikram Voleti, Samira Ebrahimi Kahou, Vincent Michalski, Tal Arbel, Chris Pal, Gaël Varoquaux, and Pascal Vincent. 2021. [Accounting for variance in machine learning benchmarks](#). In *Proceedings of Machine Learning and Systems (MLSys)*.
- Robert L. Brennan. 2001. *Generalizability Theory*. Springer.
- Ngoc Thang Bui, Guergana Savova, and Lijun Wang. 2025. [Assessing the macro and micro effects of random seeds on fine-tuning large language models](#). In *Proceedings of IJCNLP-AACL*.
- Ryan Burnell, Wout Schellaert, John Burden, and 1 others. 2023. [Rethink reporting of evaluation results in AI](#). *Science*, 380(6641):136–138.
- Arnaldo Camuffo, Alfonso Gambardella, Saeid Kazemi, Jakub Malachowski, and Abhinav Pandey. 2026. [Variance-aware LLM annotation for strategy research: Sources, diagnostics, and a protocol for reliable measurement](#). *arXiv preprint arXiv:2601.02370*.
- Jean Cardinet, Sandra Johnson, and Gianreto Pini. 2010. *Applying Generalizability Theory Using EduG*. Routledge.
- Brian E. Clauser and Kathleen M. Mazor. 1998. [Using statistical procedures to identify differentially functioning test items](#). *Educational Measurement: Issues and Practice*, 17(1):31–44.
- Mostafa Dehghani, Yi Tay, Alexey A. Gritsenko, Zhe Zhao, Neil Houlsby, Fernando Diaz, Donald Metzler, and Oriol Vinyals. 2021. [The benchmark lottery](#). *arXiv preprint arXiv:2107.07002*.
- Jesse Dodge, Gabriel Ilharco, Roy Schwartz, Ali Farhadi, Hannaneh Hajishirzi, and Noah A. Smith. 2020. [Fine-tuning pretrained language models: Weight initializations, data orders, and early stopping](#). *arXiv preprint arXiv:2002.06305*.
- Bradley Efron and Robert J. Tibshirani. 1993. *An Introduction to the Bootstrap*. Chapman & Hall, New York.
- Timo Freiesleben. 2026. [Establishing construct validity in LLM capability benchmarks requires nomological networks](#). *arXiv preprint arXiv:2603.15121*.
- Leo Gao, Jonathan Tow, Baber Abbasi, Stella Biderman, Sid Black, Anthony DiPofi, Charles Foster, Laurence Golber, Jeffrey Hsu, Alain Le Noac’h, Haonan Li, Kyle McDonell, Niklas Muennighoff, Chris Ocqueateau, Jason Phang, Laria Reynolds, Hailey Schoelkopf, Aviya Skowron, Lintang Sutawika, and 5 others. 2024. [A framework for few-shot language model evaluation](#).
- Aryo Pradipta Gema, Joshua Ong Jun Leang, Giwon Hong, Alessio Devoto, Alberto Carlo Maria Mancino, Rohit Saxena, Xuanli He, Yu Zhao, Xiaotang Du, Mohammad Reza Ghasemi Madani, Claire Barale, Robert McHardy, Joshua Harris, Jean Kaddour, Emile Van Krieken, and Pasquale Minervini. 2025. [Are we done with MMLU?](#) In *Proceedings of the Conference of the North American Chapter of the Association for Computational Linguistics (NAACL)*.
- Haiyang He and Thinking Machines Lab. 2025. [Defeating nondeterminism in LLM inference](#). Thinking Machines Lab: Connectionism.
- Peter Henderson, Riashat Islam, Philip Bachman, Joelle Pineau, Doina Precup, and David Meger. 2018. [Deep reinforcement learning that matters](#). In *Proceedings of the AAAI Conference on Artificial Intelligence*.
- Paul W. Holland and Howard Wainer, editors. 1993. *Differential Item Functioning*. Lawrence Erlbaum Associates, Hillsdale, NJ.
- Angie Hua, Kevin Tang, Chaoqun Gu, Jing Gu, Edmund Wong, and Yilin Qin. 2025. [Flaw or artifact? rethinking prompt sensitivity in evaluating LLMs](#). In *Proceedings of the Conference on Empirical Methods in Natural Language Processing (EMNLP)*.

762	Zhehan Jiang and William Skorupski. 2018. A Bayesian approach to estimating variance components within a multivariate generalizability theory framework . <i>Behavior Research Methods</i> , 50.	817
763		818
764		819
765		820
766	Ryan Othniel Kearns. 2026. Quantifying construct validity in large language model evaluations . <i>arXiv preprint arXiv:2602.15532</i> .	821
767		822
768		823
769	Douwe Kiela, Max Bartolo, Yixin Nie, Divyansh Kaushik, Atticus Geiger, Zhengxuan Wu, Bertie Vidgen, Grusha Prasad, Amanpreet Singh, Pratik Ring-shia, Zhiyi Ma, Tristan Thrush, Sebastian Riedel, Zeerak Waseem, Pontus Stenetorp, Robin Jia, Mohit Bansal, Christopher Potts, and Adina Williams. 2021. Dynabench: Rethinking benchmarking in NLP. In <i>Proceedings of the 2021 Conference of the North American Chapter of the Association for Computational Linguistics (NAACL)</i> .	824
770		
771		
772		
773		
774		
775		
776		
777		
778		
779	Alex Kipnis, Konstantinos Voudouris, Luca Schulze Buschoff, and Eric Schulz. 2025. metabench – a sparse benchmark of reasoning and knowledge in large language models . In <i>Proceedings of the International Conference on Learning Representations</i> .	825
780		826
781		827
782		828
783		
784		
785	Nitzan Kunievsky and James A. Evans. 2025. Measuring intent comprehension in LLMs: A variance decomposition framework . <i>arXiv preprint arXiv:2506.16584</i> .	829
786		830
787		
788		
789	John P. Lalor, Hao Wu, and Hong Yu. 2019. Learning latent parameters without human response patterns: Item response theory with artificial crowds . In <i>Proceedings of the Conference on Empirical Methods in Natural Language Processing and International Joint Conference on Natural Language Processing (EMNLP-IJCNLP)</i> .	831
790		832
791		833
792		834
793		835
794		836
795		
796	Guangming Li. 2023. Which method is optimal for estimating variance components and their variability in generalizability theory? Evidence from a set of unified rules for bootstrap method . <i>PLoS ONE</i> .	837
797		838
798		839
799		
800	Peiyu Li, Xiuxiu Tang, Si Chen, Ying Cheng, Ronald Metoyer, Ting Hua, and Nitesh V. Chawla. 2025. Adaptive testing for LLM evaluation: A psychometric alternative to static benchmarks . <i>arXiv preprint arXiv:2511.04689</i> .	840
801		841
802		842
803		843
804		844
805	Percy Liang, Rishi Bommasani, Tony Lee, Dimitris Tsipras, Dilara Soylu, Michihiro Yasunaga, Yian Zhang, Deepak Narayanan, Yuhuai Wu, Ananya Kumar, and 1 others. 2023. Holistic evaluation of language models. <i>arXiv preprint arXiv:2211.09110</i> .	845
806		846
807		847
808		848
809		849
810	Lovish Madaan, Samarth Bhatia, Marco Tulio Ribeiro, and Dan Roth. 2024. Quantifying variance in evaluation benchmarks . <i>arXiv preprint arXiv:2406.10229</i> .	850
811		851
812		852
813	Ian Magnusson, Noah A. Smith, and Jesse Dodge. 2023. Reproducibility in NLP: What have we learned from the checklist? In <i>Findings of the Association for Computational Linguistics (ACL)</i> .	853
814		854
815		855
816		856
	Felipe Maia Polo, Lucas Weber, Leshem Choshen, Yuekai Sun, Gongjun Xu, and Mikhail Yurochkin. 2024. tinyBenchmarks: Evaluating LLMs with fewer examples . In <i>Proceedings of the International Conference on Machine Learning (ICML)</i> .	857
		858
		859
		860
		861
		862
	William Meredith. 1993. Measurement invariance, factor analysis and factorial invariance . <i>Psychometrika</i> , 58(4):525–543.	863
		864
		865
		866
		867
		868
	Alberto Messina and Sara Scotta. 2026. Introducing background temperature to characterise hidden randomness in large language models . <i>arXiv preprint arXiv:2604.22411</i> .	
	Solomon Messing. 2026. Hidden measurement error in LLM pipelines distorts annotation, evaluation, and benchmarking . <i>arXiv preprint arXiv:2604.11581</i> .	
	Moran Mizrahi, Guy Kaplan, Dan Malkin, Rotem Dror, Dafna Shahaf, and Gabriel Stanovsky. 2024. State of what art? A call for multi-prompt LLM evaluation . <i>Transactions of the Association for Computational Linguistics</i> , 12:933–949.	
	Cilia Mushquash and Brian P. O’Connor. 2006. SPSS and SAS programs for generalizability theory analyses . <i>Behavior Research Methods</i> .	
	Yotam Perlitz, Elron Bandel, Ariel Gera, Ofir Arviv, Liat Ein-Dor, Eyal Shnarch, Noam Slonim, Michal Shmueli-Scheuer, and Leshem Choshen. 2023. Efficient benchmarking of language models . <i>arXiv preprint arXiv:2308.11696</i> .	
	Pouya Pezeshkpour and Estevam Hruschka. 2024. Large language models sensitivity to the order of options in multiple-choice questions . In <i>Findings of the North American Chapter of the Association for Computational Linguistics (NAACL)</i> .	
	Diane L. Putnick and Marc H. Bornstein. 2016. Measurement invariance conventions and reporting: The state of the art and future directions for psychological research . <i>Developmental Review</i> , 41:71–90.	
	Qi Qian, Chengsong Huang, Jingwen Xu, and 1 others. 2026. Benchmark²: Systematic evaluation of llm benchmarks . <i>arXiv preprint arXiv:2601.03986</i> .	
	Pedro Rodriguez, Joe Barrow, Alexander Hoyle, John P. Lalor, Robin Jia, and Jordan Boyd-Graber. 2021. Evaluation examples are not equally informative: How should that change NLP leaderboards? In <i>Proceedings of the Association for Computational Linguistics (ACL)</i> .	
	Angelika Romanou, Marah Ibrahim, Candace Ross, Chantal Shaib, Kerem Oktar, Samuel J. Bell, Anaelia Ovalle, Jesse Dodge, Antoine Bosselut, Koustuv Sinha, and Adina Williams. 2026. Brittlebench: Quantifying LLM robustness via prompt sensitivity . <i>arXiv preprint arXiv:2603.13285</i> .	

869	Michael Saxon, Ari Holtzman, Peter West,	Hao Zhou, Huifeng Huang, Zhe Zhao, Long Han,	922
870	William Yang Wang, and Naomi Saphra. 2024.	Haobo Wang, Ke Chen, Min Yang, Wanli Bao, Jiaxin	923
871	Benchmarks as microscopes: A call for model	Dong, Bo Xu, Conghui Zhu, Hao Cao, and Tiejun	924
872	metrology . In <i>Proceedings of the Conference on</i>	Zhao. 2026. Lost in benchmarks? rethinking large	925
873	<i>Language Modeling (COLM)</i> .	language model benchmarking with item response	926
874	Melanie Sclar, Yejin Choi, Yulia Tsvetkov, and Alane	theory . In <i>Proceedings of the AAAI Conference on</i>	927
875	Suhr. 2024. Quantifying language models’ sensitiv-	<i>Artificial Intelligence</i> . Oral presentation.	928
876	ity to spurious features in prompt design or: How I	Longyuan Zhu, Hairan Hua, Linlin Miao, and Bing	929
877	learned to start worrying about prompt formatting .	Zhao. 2026. Benchmark health index: A system-	930
878	In <i>Proceedings of the International Conference on</i>	atic framework for benchmarking the benchmarks of	931
879	<i>Learning Representations (ICLR)</i> .	LLMs . <i>arXiv preprint arXiv:2602.11674</i> .	932
880	Sanjif Shanmugavelu, Mathieu Taillefumier, Christo-	A Additional Robustness Analyses	933
881	pher Culver, Oscar Hernandez, Mark Coletti, and	We report nine robustness analyses that verify the	934
882	Ada Sedova. 2024. Impacts of floating-point non-	stability of the variance decomposition under alter-	935
883	associativity on reproducibility for HPC and deep	native analytical choices.	936
884	learning applications . In <i>SC’24 Workshops (IEEE)</i> .	GLMM complement. A generalized linear	937
885	Richard J. Shavelson and Noreen M. Webb. 1991. <i>Gen-</i>	mixed model (GLMM) with marginal logit link pro-	938
886	<i>eralizability Theory: A Primer</i> . SAGE Publications.	duces variance component rankings that perfectly	939
887	Richard J. Shavelson, Noreen M. Webb, and Glenn L.	preserve all seven grouped component orderings	940
888	Rowley. 1989. Generalizability theory . <i>American</i>	($\rho = 1.0$) relative to Henderson Method I (Jiang	941
889	<i>Psychologist</i> .	and Skorupski, 2018; Tuerlinckx et al., 2006). This	942
890	Dan Song, Won-Chan Lee, and Hong Jiao. 2025. Ex-	confirms that the linear variance decomposition is	943
891	ploring LLM autoscoring reliability in large-scale	robust to the binary outcome distribution.	944
892	writing assessments using generalizability theory .	Accuracy-range stability. The <code>item_id</code> vari-	945
893	<i>arXiv preprint arXiv:2507.19980</i> .	ance share remains stable across subsets with mean	946
894	Francis Tuerlinckx, Frank Rijmen, Geert Verbeke, and	accuracy from 0.62 to 0.72 ($r = -0.048$). This	947
895	Paul De Boeck. 2006. Statistical inference in general-	rules out distortion from binary floor or ceiling	948
896	ized linear mixed models: A review . <i>British Journal</i>	effects at observed accuracy levels.	949
897	<i>of Mathematical and Statistical Psychology</i> .	Estimator robustness. We compare Henderson	950
898	Yubo Wang, Xueguang Ma, Ge Zhang, Yuansheng Ni,	Method I with restricted maximum likelihood	951
899	Abhranil Chandra, Shiguang Guo, Weiming Ren,	(REML-EM, 500 iterations, practical convergence	952
900	Aaran Arulraj, Xuan He, Ziyang Jiang, Tianle Li, Max	$\Delta < 0.003$ percentage points). All variance com-	953
901	Ku, Kai Wang, Alex Zhuang, Rongqi Fan, Xiang	ponents differ by < 0.5 percentage points between	954
902	Yue, and Wenhui Chen. 2024. MMLU-Pro: A more	estimators. The largest discrepancy occurs on	955
903	robust and challenging multi-task language under-	the residual component (0.47pp); all substantive	956
904	standing benchmark . In <i>Advances in Neural Infor-</i>	components (item, model $\times$ item, model) differ by	957
905	<i>mation Processing Systems (NeurIPS), Datasets and</i>	< 0.3 pp. Henderson Method I is preferred for our	958
906	<i>Benchmarks Track</i> .	balanced designs because it produces unbiased esti-	959
907	Noreen M. Webb, Richard J. Shavelson, and Edward H.	mates without distributional assumptions (Brennan,	960
908	Haertel. 2006. Reliability coefficients and general-	2001 , Ch. 12).	961
909	izability theory . In C. R. Rao and Sandip Sinharay,	Probit simulation. A probit simulation matched	962
910	editors, <i>Handbook of Statistics</i> , volume 26 of <i>Hand-</i>	to our 1,296-condition design confirms that Hender-	963
911	<i>book of Statistics</i> , pages 81–124. Elsevier.	son Method I recovers variance component rank-	964
912	Haoran Ye, Jing Jin, Yuhang Xie, Xin Zhang, and Guo-	ings from dichotomized data (Spearman $\rho = 0.732$	965
913	jie Song. 2025. Large language model psychomet-	between binary and continuous rankings, exclud-	966
914	rics: A systematic review of evaluation, validation,	ing residual; <code>item_id</code> dominant in all 100 repli-	967
915	and enhancement . <i>arXiv preprint arXiv:2505.08245</i> .	cations; top-5 ordering correct in 79% of replica-	968
916	Juntao Yuan, Haoyang Li, Xingcheng Ding, Weijie Xie,	tions). Bias analysis by difficulty stratum shows	969
917	Yue-Jiao Li, Wentao Zhao, Kaiyan Wan, Jiayi Shi, Xi-	$ \text{bias} = 0.28\text{--}0.37$ for items with $p \in [0.5, 0.8]$,	970
918	aolu Hu, and Zhouhan Liu. 2025. Understanding and	where the Bernoulli variance ceiling $p(1-p)$ is	971
919	mitigating numerical sources of nondeterminism in		
920	LLM inference . In <i>Advances in Neural Information</i>		
921	<i>Processing Systems (NeurIPS)</i> . Oral presentation.		

largest. This analysis covers the observed accuracy range ($p \in [0.3, 0.9]$); extreme-difficulty items ($p < 0.3$ or $p > 0.9$) are rare in our sample (<8% of items) and may exhibit larger deviations. Henderson Method I absorbs this Bernoulli-inherent variance into the residual, explaining inflated bias near medium difficulty; systematic component ordering is robustly recovered across all replications. Binary scoring inflates residual variance by absorbing the Bernoulli $p(1-p)$ ceiling, which enters σ_{rel}^2 and increases D-study item requirements. Within a matched probit simulation, binarization increases D-study budgets by approximately 36% ($99 \rightarrow 135$ items for $G \geq 0.80$), driven by residual inflation (16% $\rightarrow$ 41%) and item variance compression (50% $\rightarrow$ 32%). Facet $\times$ item interactions—the components entering σ_{rel}^2 —are best preserved under binarization (MAPE = 12.8%, range 8.7–17.5%), while residual variance inflates as expected from Bernoulli $p(1-p)$ noise (MAPE = 156%). D-study budgets derived from binary data are therefore conservative: continuous-score designs would require equal or fewer items for the same G threshold. Null calibration (ratio = 1.0008) confirms that this conservatism does not distort relative component magnitudes.

CI method comparison. Bootstrap (1,000 item-level resamples with replacement) and delete-1 jackknife confidence intervals agree closely: width ratios range from 0.96 to 1.04 across all variance components. Bootstrap CIs capture item-sampling variability conditional on these 8 specific models; model-level stability is assessed separately via LOMO (below). The agreement indicates that the dependence structure introduced by item-level bootstrap resampling (which creates ~ 127 unique items per 200-item resample) does not materially distort interval estimates. Stratified resampling by subject (10 subjects $\times$ 20 items) produces CIs within 4% of simple item-level bootstrap, confirming that the within-subject dependence structure does not materially affect interval estimates.

Facet specification sensitivity. Treating temperature and ordering as fixed (rather than random) facets changes G_{item} by only 0.37 percentage points (all-random: 0.820 vs. fixed temperature+ordering: 0.824). The small difference reflects the negligible temperature $\times$ item (0.65%) and ordering $\times$ item (0.54%) variance shares. The random-facet assumption is therefore robust to facet classification.

Leave-one-model-out (LOMO). Across all 40 leave-one-model-out subsets (8 models $\times$ 5 benchmarks), G_{item} ranges from 0.833 to 0.919, with CV $\leq 1.25\%$ per benchmark. No single model removal changes G by more than 0.030 (GSM8K, removing Gemma-2-9B). Item-related variance rankings are preserved in all 40 subsets. Table 5 reports per-benchmark stability. This supports treating models as exchangeable within the target universe of 7–9B instruction-tuned models (§3.6).

Table 5: LOMO stability across five benchmarks (8 models $\times$ 5 benchmarks = 40 subsets).

Benchmark	G_{full}	CV (%)	Max $ \Delta G $	Most impactful
MMLU	0.896	0.49	0.019	InternLM2.5-7B
ARC	0.892	0.94	0.026	Llama-3.1-8B
HellaSwag	0.872	1.09	0.028	Llama-3.1-8B
GSM8K	0.863	1.25	0.030	Gemma-2-9B
MATH	0.888	1.20	0.029	InternLM2.5-7B

Inter-model agreement. We computed pairwise Cohen’s κ and Fleiss’ κ across all model pairs on each benchmark. Mean pairwise κ ranges from 0.326 (MATH) to 0.461 (MMLU); Fleiss’ κ ranges from 0.309 (GSM8K) to 0.456 (MMLU). HellaSwag uses six models ($\kappa = 0.332$) due to missing OLMo and Yi data. The fair-to-moderate agreement independently confirms that the model $\times$ item interaction (30–36% of variance) reflects genuine differential item functioning: models disagree substantially on which items are difficult, consistent with prior DIF findings in educational measurement (Holland and Wainer, 1993; Clauser and Mazor, 1998).

Scoring metric sensitivity. To probe metric conditionality, we compared binary correctness with step-level partial credit on GSM8K using the full 8-model crossed design (200 items, 2,304 conditions per item). Under partial credit, G_{item} drops from 0.863 to 0.768, with residual variance rising from 39.9% to 64.9% and item variance falling from 16.9% to 10.3%. All eight models individually show $G_{\text{partial}} < G_{\text{binary}}$ (mean $\Delta G = -0.135$, range -0.025 to -0.309), with ΔG magnitude moderated by model accuracy: higher-accuracy models have fewer incorrect responses whose scores shift under partial credit (e.g., OLMo at 70.1% accuracy shows $\Delta G = -0.025$, while Llama at 49.0% shows $\Delta G = -0.309$). The mechanism is clear: 48.1% of incorrect responses have numerically correct arithmetic steps but incorrect problem setup, receiving partial score = 1.0 de-

spite binary score = 0—this compresses item discrimination and inflates residual variance. Binary scoring amplifies prompt sensitivity by collapsing reasoning quality into pass/fail, confirming that all findings in this paper are metric-conditional.

Facet specification sensitivity: random vs. fixed model. Table 6 reports per-benchmark generalizability coefficients under random and fixed model specifications. The consistent ΔG of 0.077–0.113 confirms that random-model budgets are strictly more conservative across all benchmarks.

Table 6: Generalizability coefficients under random vs. fixed model facet specification across five benchmarks.

Benchmark	G_{random}	G_{fixed}	ΔG
MMLU	0.896	0.984	+0.088
ARC	0.892	0.990	+0.098
HellaSwag	0.872	0.984	+0.113
GSM8K	0.863	0.951	+0.089
MATH	0.888	0.965	+0.077
Mean	0.882	0.975	+0.093

Null calibration. We generate 100 Bernoulli-distributed datasets with item difficulties matched to the empirical distribution but no systematic facet effects. The observed-to-null variance ratio is 1.0008, confirming that Henderson Method I does not introduce systematic bias when applied to binary outcomes. The null simulation also verifies that the item variance share (58.4%) is not an artifact of the Bernoulli $p(1-p)$ ceiling: the null produces item variance of 0.08%, three orders of magnitude below the observed value. Across all 100 replications, the same three components (item, model $\times$ item, residual) emerge as the top three (set consistency 100%), with Kendall’s $\tau = 0.81$ between binary and continuous rankings.

Reduced crossing validation. To verify that the reduced crossing in exp-002 (BF16 only $\times$ 2 temperatures) does not introduce confounding, we extracted the corresponding subset from the fully crossed exp-001 design. The generalizability coefficient differs by only 0.0008 ($G = 0.9361$ vs. 0.9369), and all non-precision variance components differ by less than 2 percentage points. The 10pp increase in item variance share (58.4% $\rightarrow$ 68.5%) is fully explained by the removal of precision-related components ($\sim 4.5\%$ total), which redistributes across remaining sources. This confirms that the exp-002 design preserves the variance structure of the full crossing.

B Cross-Model D-Study Projections

Table 7 extends the single-model D-study (Table 4 in main text) to per-benchmark cross-model projections with bootstrap confidence intervals (1,000 item-level resamples, 8 models, 2,304 conditions per item). The efficiency ratio is format-associated: MC benchmarks show 11–13 $\times$ (mean 12.0 $\times$ [11.6, 12.5]), while FF benchmarks show 21–22 $\times$ (mean 21.5 $\times$ [20.4, 22.8]).

Table 7: Cross-model D-study reproducibility budgets with 95% bootstrap CIs. $n_{G \geq 0.80}$: items needed. Multi = multi-facet (2,304 cond.); Total = items $\times$ 2,304. Single = single-condition (total evals = item count).

	Multi-facet (2,304 cond.)			Single-cond.		
	Items	95% CI	Total	Items	95% CI	Ratio [95% CI]
<i>Multiple-choice</i>						
MMLU	93	[76, 116]	214K	1,216	[1004, 1491]	13.1 [12.5, 13.7]
ARC	98	[71, 139]	226K	1,078	[793, 1512]	11.0 [10.7, 11.4]
Hella.	118	[93, 151]	272K	1,403	[1120, 1790]	11.9 [11.5, 12.3]
<i>Free-form</i>						
GSM8K	128	[103, 162]	295K	2,768	[2230, 3531]	21.6 [20.7, 23.0]
MATH	101	[86, 122]	233K	2,160	[1853, 2631]	21.4 [20.2, 22.5]

The higher item requirements relative to the single-model analysis (55 items in Table 4) reflect the additional model $\times$ item interaction variance that enters σ_{rel}^2 when models are crossed as a facet. The format-associated ratio pattern—FF benchmarks showing $\sim 2\times$ the MC ratio—is consistent with the distributed variance structure observed in §4.3, where FF benchmarks allocate more variance to model and model $\times$ prompt components.

Internal consistency check. Table 8 reports consistency checks from four exp-001 models (DeepSeek, Llama, Mistral, Qwen3) to the full eight-model design across all five benchmarks. MC benchmarks show $|\Delta G| \leq 0.029$ ($< 3.5\%$ error), while FF benchmarks show larger deviations ($|\Delta G| \leq 0.065$, $\sim 7\%$) due to limited model diversity in the four-model subset. FF predictions consistently underestimate actual G (conservative); MC predictions show mixed directions with small deviations ($|\Delta G| \leq 0.029$).

Budget sensitivity to estimation error. Figure 4 shows the effect of $|\Delta G| = 0.035$ estimation error on item budgets across all five benchmarks. The percentage increase is remarkably stable (23.0–24.5%, mean 23.2%), because the relative budget adjustment $\Delta n/n$ depends on the G -formula ratio rather than on absolute variance components. Absolute Δn ranges from 13 (ARC) to 67 (GSM8K), scaling with the baseline budget.

Table 8: Internal consistency: four-model subset vs. eight-model G_{item} . FF subsets underestimate actual G (conservative); MC shows mixed directions with small deviations.

Benchmark	Predicted G	Actual G	$ \Delta G $	Error%
<i>Multiple-choice</i>				
MMLU	0.888	0.896	0.009	0.96
ARC	0.904	0.892	0.012	1.38
HellaSwag	0.901	0.872	0.029	3.33
<i>Free-form</i>				
GSM8K	0.803	0.863	0.059	6.88
MATH	0.823	0.888	0.065	7.30
Mean	—	—	0.035	3.97




Figure 4: Effect of $|\Delta G| = 0.035$ estimation error on item budgets across five benchmarks. The $\sim 23\%$ budget increase is consistent across all benchmarks, indicating D-study budgets are robust to estimation uncertainty.

Cross-benchmark G stability. G_{item} values cluster in a narrow band across all five benchmarks (0.863–0.896, CV = 1.4%), despite qualitatively different variance structures: MC benchmarks concentrate 68–72% in item-related components while GSM8K disperses variance across item (17%), model (11%), and interaction (15%) sources (Table 3). This stability implies that D-study budget recommendations calibrated on one benchmark provide reasonable order-of-magnitude guidance for others. However, the 0.033 range in G_{item} translates to a 71–128 item range for multi-facet $G \geq 0.80$ (Table 7), and the variance compositions differ substantially—benchmark-specific G-studies remain necessary to identify dominant variance sources and optimize condition selection.

C Intermediate D-Study Designs

Table 9 shows how reliability improves as evaluation designs increase from single-condition to full multi-facet crossing. The steepest gain occurs at the “minimal upgrade” level (2 prompts): G_{200} jumps from 0.592 to 0.692 at $2\times$ cost. A minimal sufficient design for $G \geq 0.80$ is 3 prompts $\times$ 3 seeds (9 conditions, $G_{200} = 0.818$, $n_{G \geq 0.80} = 178$). Beyond the “standard” level (32 conditions), adding orderings provides negligible benefit (ordering $\times$ item $< 1\%$), confirming that

prompt and seed diversity are the primary reliability drivers.

Table 9: Intermediate D-study designs (single-model, MMLU). Designs are ordered by increasing conditions per item; G_{200} and $n_{G \geq 0.80}$ are D-study projections from exp-001 variance components.

Design	Conditions	G_{200}	$n_{G \geq .80}$
Single-condition	1	.592	552
Minimal upgrade (2p)	2	.692	357
Easy upgrade (2p $\times$ 2s)	4	.766	245
Moderate (3p $\times$ 3s)	9	.818	178
Standard (4p $\times$ 4s $\times$ 2t)	32	.869	121
Enhanced (+4 orderings)	72	.871	119
Full multi-facet	1,296	.936	55

D Format-Associated Variance: Per-Benchmark Details

This section supplements the compressed format-associated analysis in §4.3 with per-benchmark details.

MC benchmark exceptions. HellaSwag is the sole MC case where model $\times$ item slightly exceeds item (36.0% vs. 34.3%), though their sum (70.3%) remains in the MC range (68–72%). Among MC benchmarks, ARC-Challenge shows an anomalously high model $\times$ prompt interaction (8.0% vs. $< 2\%$ for MMLU and HellaSwag), suggesting that science-domain items are differentially sensitive to option formatting across model families.

FF benchmark patterns. Free-form benchmarks differ markedly from MC. On GSM8K, variance distributes across item (16.9%), model $\times$ item (14.6%), model (10.7%), and model $\times$ prompt (8.5%); prompt main effect is 5.4%. MATH shows an intermediate pattern: item leads at 24.0% with model $\times$ item at 17.3%, but residual variance is substantially larger (43.9%). Item-level regression reveals that FF residual variance is predominantly explained by the Bernoulli variance ceiling $p(1-p)$ ($R^2 = 0.47\text{--}0.86$; Appendix G): medium-difficulty items maximize binary outcome variance regardless of facet structure, so the elevated FF residuals reflect measurement resolution rather than unmodeled substantive sources.

Difficulty dispersion as variance predictor. Difficulty dispersion ($\text{std}(p_i)$) strongly predicts item variance share ($\rho = 0.82$, $n = 5$ benchmarks; Appendix E; reasoning-heavy subjects show G_{item} as low as 0.614, Appendix F). Format constrains the noise floor, but difficulty heterogeneity determines variance magnitude within it.

Item-level bootstrap stability. Bootstrap resampling (1,000 item resamples per benchmark) confirms that the MC–FF separation is robust to item selection: 95% CIs for item-related variance (item + model $\times$ item) are [65.2, 75.0]% (MC) and [28.0, 44.2]% (FF), with no overlap (minimum gap 21.0 percentage points).

E Cross-Benchmark Difficulty Heterogeneity

Figure 5 plots item difficulty heterogeneity ($\text{std}(p_i)$ across models) against item_id variance share for all five benchmarks. MC benchmarks cluster at higher dispersion (0.27–0.31) and higher item variance (34–38%), while FF benchmarks occupy the lower-left region (0.22–0.23, 17–24%). The cross-benchmark Spearman $\rho = 0.82$ ($p = 0.089$, $n = 5$) is consistent in direction with the within-MMLU correlation ($\rho = 0.806$, $p = 0.005$, $n = 10$ subjects); these are independent analyses at different granularities (benchmarks vs. subjects within MMLU), providing multi-scale support for the item difficulty heterogeneity mechanism discussed in §4.3.




Figure 5: Item difficulty heterogeneity ($\text{std}(p_i)$, computed over available models per benchmark, 6–8 models) versus item_id variance share (as reported in Table 3) across five benchmarks. MC benchmarks (filled) exhibit higher difficulty dispersion and item variance than FF benchmarks (open). Spearman $\rho = 0.82$, $p = 0.089$.

F MMLU Subject-Level Variance Decomposition

Table 10 reports per-subject variance decomposition for 10 MMLU subjects (20 items each, 8 models, 2,304 conditions per item). G_{item} ranges from 0.614 (*high_school_mathematics*) to 0.930 (*international_law*), with mean 0.843. Knowledge-intensive

subjects (*international_law*, *clinical_knowledge*, *world_religions*) show $\text{item_id} > 40\%$, while reasoning-heavy subjects (*high_school_mathematics*, *abstract_algebra*) show $\text{model} \times \text{item} > \text{item_id}$, consistent with models differing more in reasoning strategies than in factual recall. The *high_school_mathematics* item_id of 10.4% referenced in §4.3 is the lowest across all subjects, below even GSM8K (16.9%).

Table 10: Subject-level variance decomposition within MMLU (20 items per subject, 8 models, 2,304 conditions per item). Sorted by G_{item} descending.

Subject	Acc.	item%	mod. $\times$ it.%	G_{item}
<i>international_law</i>	.717	44.4	24.1	.930
<i>clinical_knowledge</i>	.581	47.5	30.0	.923
<i>world_religions</i>	.826	42.2	27.1	.910
<i>computer_security</i>	.740	37.0	24.9	.908
<i>college_physics</i>	.436	31.6	30.1	.876
<i>philosophy</i>	.782	33.0	36.1	.868
<i>abstract_algebra</i>	.457	24.1	32.8	.822
<i>us_foreign_policy</i>	.864	20.4	36.9	.802
<i>professional_medicine</i>	.671	19.9	41.7	.778
<i>high_school_math</i>	.385	10.4	41.9	.614
Mean	—	31.1	32.6	.843

G FF Residual Variance Analysis

To investigate whether the elevated residual variance in free-form benchmarks (GSM8K 39.9%, MATH 43.9%) reflects unmodeled substantive sources or measurement properties of binary scoring, we regress item-level residual variance on item characteristics across all five benchmarks (Table 11).

The dominant predictor is the quadratic difficulty effect ($p < 10^{-5}$ on all benchmarks): items near $p = 0.5$ show maximal residual variance, while items near the floor or ceiling show minimal residual. This inverted-U pattern directly mirrors the Bernoulli variance ceiling $p(1-p)$. Item content features (number of reasoning steps, answer length) are not significant predictors for GSM8K ($p > 0.38$), confirming that the difficulty–residual relationship is statistical rather than content-driven.

FF benchmarks show substantially higher R^2 (mean 0.66) than MC benchmarks (mean 0.25), consistent with FF items spanning a wider difficulty range. The MC–FF contrast in residual magnitude is therefore a mathematical consequence of the difficulty distribution under binary scoring, not evidence of additional unmodeled variance sources in free-form evaluation.

Table 11: Item-level residual variance regression. Difficulty² = quadratic (inverted-U) term reflecting the Bernoulli $p(1-p)$ ceiling. GSM8K additionally tests n_steps and answer_length (both $p > 0.38$).

Benchmark	R^2 (full)	R^2 (diff. only)	Diff. ² p
<i>Multiple-choice</i>			
MMLU	0.377	0.347	<0.001
ARC	0.221	0.204	<0.001
HellaSwag	0.156	0.140	<0.001
<i>Free-form</i>			
GSM8K	0.468	0.449	<0.001
MATH	0.856	0.845	<0.001

H Model and Prompt Details

Models. Table 12 lists the eight models with their HuggingFace identifiers.

Table 12: Models used in all experiments.

Short name	HuggingFace identifier	Params
Llama-3.1-8B	meta-llama/Llama-3.1-8B-Instruct	8B
Gemma-2-9B	google/gemma-2-9b-it	9B
Mistral-7B-v0.3	mistralai/Mistral-7B-Instruct-v0.3	7B
Qwen3-8B	Qwen/Qwen3-8B	8B
Yi-1.5-9B	01-ai/Yi-1.5-9B-Chat	9B
OLMo-2-7B	allenai/OLMo-2-1124-7B-Instruct	7B
InternLM2.5-7B	internlm/internlm2_5-7b-chat	7B
DeepSeek-7B	deepseek-ai/deepseek-llm-7b-chat	7B

Prompt templates. Our six prompt templates follow a surface perturbation design inspired by BrittleBench (Romanou et al., 2026). For multiple-choice benchmarks, templates vary along three dimensions: option layout (lettered vs. numbered), instruction wording (direct vs. contextual), and answer prompt cue. For free-form benchmarks, templates vary chain-of-thought elicitation strategy. All six templates are semantically equivalent—they request the same task with different surface forms—to isolate formatting effects from content effects. ARC-Challenge and HellaSwag use the same six surface patterns as MMLU with domain-appropriate system instructions; MATH mirrors the GSM8K pattern with `\boxed{}` answer formatting.

MMLU templates (representative MC format):

1. The following are multiple choice questions (with answers) about {subject}.
... D. ...
Answer:
2. Answer the following {subject} questions by selecting the correct option.
Question:

{question}
(A) ... (B) ...
(C) ... (D) ...
Correct
answer:

3. Below are {subject} multiple-choice problems. Pick the right letter.
Question:
{question}
(A) ... (B) ...
(C) ... (D) ...
Answer:
4. Answer the following {subject} question by selecting A, B, C, or D.
{question}
- A. ... - B. ... - C. ... - D. ...
Answer:
5. Below is a {subject} exam question with four possible answers. Choose the correct one.
{question}
Option A: ... Option B: ... Option C: ... Option D: ...
Correct:
6. Question ({subject}):
{question}
(A) ... (B) ... (C) ... (D) ...
Your answer:

GSM8K templates (representative FF format):

1. Solve the following math problem step by step.
Question:
{question}
Answer: Let's think step by step.
2. Answer the following math word problem. Show your work.
Question:
{question}
Solution:
3. Solve each problem below. Write your reasoning, then give the final answer after ####.
Problem:
{question}
Reasoning:
4. You are a math tutor. Solve the problem and show your steps.
{question}
Steps:
5. Work through this math question carefully.
Question:
{question}
Work:
6. Math problem:
{question}
Let's solve this step by step.

I Use of AI Assistants

We used Claude (Anthropic) as an AI coding and writing assistant throughout this work. Specifically:

- **Code:** drafting experiment scripts, data-processing pipelines, and analysis utilities; debugging and refactoring.
- **Writing:** drafting and revising paper sections, polishing prose, and reformatting L^AT_EX.
- **Literature:** summarizing related papers to inform related-work section.