

Reproducibility Budgets for Open-Weight LLM Evaluation: A Generalizability Theory Study of 7–9B Instruction-Tuned Models

Anonymous ACL submission

Abstract

Large Language Model (LLM) evaluation results vary across reproductions, but variance sources interact, and their relative importance shifts with the benchmark format. We apply Generalizability theory (G-theory) in a fully crossed design with up to seven facets across open-weight instruction-tuned models and five benchmarks. Under binary correctness scoring, a recommended 3-prompt $\times$ 3-seed design achieves $G \geq 0.80$ with 178 items—a $3.1\times$ item reduction over single-condition designs at the cost of $2.9\times$ more total evaluations—a tradeoff that favors multi-facet designs when item construction is costly. Variance decomposition reveals that models differ primarily in which items they find difficult rather than in overall ability, and that the variance structure shifts qualitatively between multiple-choice and free-form benchmarks. Preliminary validation across model scales and architecture families confirms structural stability. We release an open-source G-theory pipeline and D-study budget calculator for evaluation design optimization.

1 Introduction

In our analysis of five benchmarks, item selection drives 34–38% of evaluation variance on multiple-choice tasks while prompt wording is negligible ($<3\%$). On the free-form reasoning task GSM8K, the structure shifts: model $\times$ prompt interaction alone accounts for approximately 9%, while item selection drops to 17%. These numbers, drawn from a factorial analysis of approximately 2.3 million evaluation records, reveal that the variance structure of LLM evaluation is not a fixed property of the pipeline but shifts qualitatively with the response format. Current evaluation practice treats variance sources as independent nuisance factors and controls them one at a time: one study perturbs prompts, another varies precision, a third changes

item subsets. This piecemeal approach cannot capture interactions between sources, nor can it predict how the relative importance of each source changes across evaluation settings. Leaderboard rankings built on single-condition evaluations are therefore sensitive to methodological choices that remain unreported and uncontrolled (Rodriguez et al., 2021; Dehghani et al., 2021). No study has simultaneously crossed infrastructure and methodology facets in a unified factorial design across benchmarks to decompose these sources, quantify their interactions, and compare the resulting variance structure.

Existing studies examine individual variance sources in isolation. Sclar et al. (2024) and BrittleBench (Romanou et al., 2026) documented 50–76% accuracy swings from prompt formatting; Yuan et al. (2025) showed up to 9% variation from numerical precision; Pezeshkpour and Hruschka (2024), Dodge et al. (2020), and Henderson et al. (2018) documented ordering and hyperparameter sensitivity—each holding other factors fixed. Messing (2026) applied G-theory to LLM evaluation across five methodology facets on a single benchmark (judge identity: 43.8%), but excluded infrastructure facets and cross-benchmark comparison. Madaan et al. (2024) proposed per-factor metrics without crossing factors. IRT approaches (Maia Polo et al., 2024; Lalor et al., 2019; Zhou et al., 2026) optimize item selection but ignore prompt, seed, and ordering facets. Bouthillier et al. (2021) established ANOVA-based variance decomposition for ML reproducibility. Three gaps remain: no study has (a) crossed infrastructure and methodology facets in a unified factorial design, (b) compared variance structures across response formats, or (c) translated variance components into design-dependent reproducibility budgets.

Through a factorial G-theory study (Shavelson and Webb, 1991; Brennan, 2001) crossing up to seven facets across eight open-weight 7–9B param-




Figure 1: Overview of the G-theory evaluation pipeline. A factorial design crossing seven facets across eight models and five benchmarks produces 2.3M binary records. G-study variance decomposition via Henderson Method I reveals that model×item interaction accounts for up to 36% of variance. D-study projections translate variance components into design-dependent item budgets: a 3-prompt × 3-seed design achieves $G \geq 0.80$ with 178 items (3.1× reduction). Two contributions emerge: reproducibility budgets (C1) and format-associated variance patterns (C2).

eter models and five benchmarks (MMLU, ARC-Challenge, HellaSwag, GSM8K, MATH), totaling approximately 2.3 million records, we find that evaluation variance has a stable but format-dependent internal structure. D-study projections show that a recommended 3-prompt × 3-seed design achieves $G \geq 0.80$ with 178 items (1,602 evaluations per model), a 3.1× reduction from 552 single-condition items; the full 1,296-condition design further reduces to 55 items as a theoretical ceiling (Figure 2). Model×item interaction (up to 36%) indicates that models differ primarily in *which items* they find difficult (Pezeshkpour and Hruschka, 2024; Maia Polo et al., 2024).

Our contributions:

- Design-dependent reproducibility budgets.** A recommended 3-prompt × 3-seed design achieves $G \geq 0.80$ with 178 items (1,602 evaluations per model), a 3.1× reduction over single-condition designs (552 items); the full 1,296-condition design reduces to 55 items (71,280 evaluations). Under binary scoring, BF16 and FP16 require no separate control (<0.1% interaction). IRT item selection and G-theory protocol optimization compose: IRT-selected 116 items yield $G = 0.457$ single-condition but $G = 0.895$ multi-facet (Table 6).
- Format-associated variance patterns (exploratory).** MC benchmarks concentrate 68–72% of variance in item-related components; FF benchmarks disperse to 31–41% (item-level $d = 1.73$; benchmark-level confirma-




Figure 2: D-study projections for MMLU (single-model baseline). A recommended 3-prompt × 3-seed design reduces the budget from 552 to 178 items for $G \geq 0.80$; the full multi-facet design reduces further to 55 items.

tion limited by sample size). This hypothesis-generating observation warrants confirmation with larger benchmark samples.

We release an open-source G-theory pipeline and D-study budget calculator for evaluation design optimization.

2 Related Work

Variance Sources in LLM Evaluation. Bouthillier et al. (2021) introduced ANOVA-based variance decomposition for ML reproducibility, establishing that design choices account for a substantial share of benchmark variance. Subsequent work isolated individual sources: Yuan et al. (2025) documented up to 9% accuracy variation from numerical precision; Sclar et al. (2024) and BrittleBench (Romanou et al., 2026) found 50–76% accuracy swings from prompt

formatting; Pezeshkpour and Hruschka (2024) and Bui et al. (2025) documented ordering and seed sensitivity. Messina and Scotta (2026) formalized CUDA nondeterminism from floating-point non-associativity. Each study holds other factors fixed, precluding interaction estimation. Madaan et al. (2024) proposed per-factor metrics without crossing factors; holistic evaluation frameworks (Liang et al., 2023; Gao et al., 2024) reduce implementation variance but do not decompose it. Our work crosses seven facets simultaneously, enabling interaction estimation and cross-format comparison.

IRT and Efficient Benchmarking. IRT models item difficulty and discrimination to optimize benchmark efficiency (Lalor et al., 2019; Rodriguez et al., 2021). Maia Polo et al. (2024) and Zhou et al. (2026) developed IRT-based subsampling for efficient evaluation; Perlitz et al. (2023) and Kipnis et al. (2025) explored related compression strategies. These approaches optimize item selection but do not model variance from prompt, seed, or ordering facets. IRT-selected items under a single methodological condition can therefore underestimate required item counts—our results show IRT’s 116 items yield $G = 0.457$ single-condition but $G = 0.895$ multi-facet, demonstrating that item selection and protocol optimization are orthogonal.

G-Theory and Positioning. G-theory (Shavelson and Webb, 1991; Brennan, 2001) provides the canonical framework for multi-facet reliability analysis, with extensions to binary data (Mushquash and O’Connor, 2006; Jiang and Skorupski, 2018). We operationalize differential item functioning (Holland and Wainer, 1993; Clauser and Mazor, 1998) as model×item interaction. Recent psychometric perspectives on LLM evaluation address construct validity (Kearns, 2026) and measurement properties (Ye et al., 2025). Closest to our work, Messing (2026) applied G-theory to LLM evaluation, decomposing five methodology facets on a single benchmark (judge identity: 43.8%). We extend this in three directions: (i) replacing judge identity with infrastructure facets and additional methodology facets; (ii) crossing facets across five benchmarks for cross-format comparison; and (iii) translating variance components into reproducibility budgets via D-study projections. Table 1 summarizes these distinctions.

Table 1: Positioning relative to prior multi-facet variance studies. “–” = not addressed.

	Bouthillier et al. (2021)	Messing (2026)	Ours
Facets crossed	4	5	7
Design	Factorial	Factorial	Factorial
Method	ANOVA	G-theory	G-theory
D-study budgets	–	–	✓
Models	ML models	6 LLMs	8 LLMs
Benchmarks	3	1	5
Cross-format	–	–	MC+FF

3 Method and Experimental Design

3.1 Generalizability Theory

Generalizability theory (G-theory) treats each measurement as one observation from a universe of admissible conditions defined by the crossed facets of an evaluation design (Shavelson and Webb, 1991; Brennan, 2001). The *object of measurement* is the entity whose consistency is of interest. In our setting, test items (p) serve as the object of measurement, and the goal is to determine how reliably item-level scores generalize across the methodological conditions under which they are collected.

A *G-study* decomposes total observed variance into components σ_x^2 for each facet x and their interactions. Each component quantifies how much variability that source contributes to the overall measurement. Large interaction components (e.g., σ_{pM}^2 for item-by-model) indicate that facets do not affect items uniformly but instead create differential difficulty profiles. We estimate variance components via Henderson Method I, an unbiased ANOVA-based estimator for balanced designs (Brennan, 2001, Ch. 3), with 1,000-resample bootstrap confidence intervals (Efron and Tibshirani, 1993) (item-level resampling with replacement, preserving the fully-crossed structure across all other facets) (Li, 2023).

A *D-study* projects the generalizability coefficient under hypothetical designs with different numbers of conditions per facet:

$$G_{\text{item}} = \frac{\sigma_p^2}{\sigma_p^2 + \sigma_{\text{rel}}^2} \quad (1)$$

where $\sigma_{\text{rel}}^2 = \sum_i \sigma_{p \times f_i}^2 / n_{f_i} + \sigma_{\text{residual}}^2 / \prod_i n_{f_i}$ aggregates all item-by-facet interaction variances, each divided by the number of conditions for the corresponding facets (Brennan, 2001, Ch. 4). A D-study projects reliability for a test averaging over n_p items via

$$G(n_p) = \frac{n_p \sigma_p^2}{n_p \sigma_p^2 + \sigma_{\text{rel}}^2} \quad (2)$$

enabling independent verification of all budget projections in Table 6. This prospective capability distinguishes G-theory from classical ANOVA: rather than only attributing variance post hoc, a D-study enables optimization of evaluation designs *before* data collection (Shavelson et al., 1989). We define *item-related variance* as the sum of the item main effect (σ_p^2) and the model $\times$ item interaction (σ_{pM}^2). The latter captures differential item functioning (DIF): items whose difficulty ranking changes across models (Brennan, 2001, Ch. 4). Because DIF is operationally an item-level property—it reflects model-specific item difficulty profiles rather than a global model ability shift—aggregating $\sigma_p^2 + \sigma_{pM}^2$ provides a unified measure of item-driven variance. We report unbundled components throughout to ensure transparency.¹

3.2 Experimental Design: MMLU

We design a fully crossed factorial experiment (exp-001) that simultaneously varies seven facets of LLM evaluation on MMLU. Table 2 summarizes the complete design matrix. Figure 1 illustrates the complete pipeline from factorial design through variance decomposition to budget projection.

Facets. Items (p): 200 MMLU questions selected via purposive stratified sampling across 10 of MMLU’s 57 subjects (20 per subject, deterministic selection with arbitrarily chosen seed 42). The 10 subjects span four disciplinary domains—STEM: abstract algebra, college physics, high school mathematics, computer security; humanities: world religions, philosophy; social sciences: US foreign policy, international law; professional: professional medicine, clinical knowledge—ensuring that variance estimates are not dominated by a single knowledge domain. Resampling stability (CV = 3.1%) was evaluated within these 10 subjects; cross-subject generalization remains untested. Models (M): eight open-weight 7–9B parameter instruction-tuned models spanning eight architecture families—DeepSeek-7B, Gemma-2-9B, InternLM2.5-7B, Llama-3.1-8B-Instruct, Mistral-7B-v0.3-Instruct, OLMo-2-7B, Qwen3-8B, and Yi-1.5-9B (full identifiers in Appendix H)—treated as a quasi-random facet (§3.6). Numerical precision (π): BF16, FP16, and FP32, all evaluated with `enforce_eager=True` for fair compari-

¹Notation: p = item, M = model, π = precision, δ = prompt, s = seed, τ = temperature, o = ordering, σ_x^2 = variance component for source x .

Table 2: Fully crossed G-theory design. Exp-001 evaluates MMLU across all facet levels; exp-002 extends to five benchmarks with reduced crossing. [†]BF16 balanced subset (2 temperatures: 0.0, 0.7) for cross-model analysis; the full single-model design uses all 3 levels.

Facet	Levels	exp-001	exp-002
Item (p)	Stratified sample	200	200 $\times$ 5
Model (M)	8 architecture families	8	8
Precision (π)	BF16 / FP16 / FP32	3	1
Prompt (δ)	Paraphrases	6	6
Seed (s)	Fixed set	6	6
Temperature (τ)	0.0 / 0.3 / 0.7	3	2
Ordering (o)	Demo permutations	4	4
Total records		460,800 [†]	2,304,000

son across precision levels. Prompt template (δ): six surface perturbations of the evaluation instruction (Sclar et al., 2024), varying option layout (lettered vs. numbered), instruction wording (direct vs. contextual), and answer prompt cue for MC benchmarks; for FF benchmarks, templates vary chain-of-thought elicitation strategy (full template text in Appendix H). Random seed (s): six fixed values (42, 123, 456, 789, 1024, 2048). Sampling temperature (τ): three levels (0.0, 0.3, 0.7) spanning deterministic to moderately stochastic decoding. Answer ordering (o): four permutations of few-shot demonstration examples (for MMLU; structurally inert for other benchmarks, see §3.3) (Pezeshkpour and Hruschka, 2024).

Sample sizes. The cross-model BF16 balanced subset yields 460,800 records (8 models $\times$ 200 items $\times$ 6 prompts $\times$ 6 seeds $\times$ 2 temperatures $\times$ 4 orderings). A single-model analysis (Llama) spans all three precisions, yielding 259,200 binary observations across 1,296 fully balanced conditions. The BF16 subset serves as the primary cross-model sample because BF16 is the standard inference precision for these architectures.

3.3 Cross-Benchmark Extension

To test whether the variance structure observed on MMLU generalizes across evaluation formats, exp-002 extends the design to five benchmarks: MMLU, ARC-Challenge, and HellaSwag (multiple-choice), and GSM8K and MATH (free-form generation). The design crosses 8 models $\times$ BF16 $\times$ 2 temperatures (0.0, 0.7) $\times$ 6 prompts $\times$ 6 seeds $\times$ 4 orderings $\times$ 200 items per benchmark, totaling 2,304,000 records. This reduced crossing fixes precision at BF16 and uses two temperature levels to concentrate statistical power on format-associated differences. The ordering facet operates only on

MMLU, where it selects among four permutations of the few-shot demonstration examples. For all other benchmarks, the prompt construction does not use the ordering parameter; ordering is structurally inert ($\sigma_{\text{ordering}}^2 = 0$ by construction) and is retained solely for design balance. The two temperature levels bracket the most common evaluation configurations: deterministic ($\tau = 0.0$) and a moderate stochastic setting ($\tau = 0.7$) representative of practical deployment. A subset analysis confirms that this reduced crossing preserves the variance structure of the full design ($\Delta G = 0.0008$; Appendix A).

3.4 Evaluation Protocol

All evaluations use generation-based scoring (free-text response with answer extraction) rather than log-likelihood ranking, following Hua et al. (2025). Binary correctness (1/0) is the primary metric; text-exact-match is secondary. Inference uses vLLM 0.9.2 without batch-invariant kernels, preserving CUDA-level nondeterminism as a measurable variance source (Messina and Scotta, 2026; He and Thinking Machines Lab, 2025). Few-shot protocols: 5-shot MMLU/ARC, 4-shot MATH, 8-shot GSM8K, 10-shot HellaSwag.

3.5 Binary Data Assumptions

Standard G-theory assumes continuous, approximately normal outcome data (Shavelson and Webb, 1991). Our binary correctness scores (0/1) violate this assumption. We establish robustness through simulation and estimator properties, with two additional analyses (GLMM complement, accuracy-range stability) in Appendix A.

Henderson Method I produces unbiased estimates for balanced designs regardless of outcome distribution (Brennan, 2001, Ch. 12), and REML estimates differ by $<0.5\text{pp}$ on all components. A probit simulation, GLMM complement ($\rho = 1.0$), and accuracy-range analysis all confirm robustness (Appendix A).

All findings in this paper are conditioned on binary correctness scoring. This measurement resolution collapses continuous generation quality into a pass/fail outcome, which may amplify item-related variance by reducing within-item response variability and mask finer-grained variance sources (e.g., partial credit, reasoning quality) that would emerge under richer scoring rubrics. The text-exact-match comparison in §4.1 ($G = 0.007$ vs. 0.936) illustrates how the scoring metric can fundamentally

reshape variance structure. A partial-credit analysis on GSM8K (8 models, step-level scoring) confirms this conditionality: G_{item} drops from 0.86 to 0.77, with all eight models showing $G_{\text{partial}} < G_{\text{binary}}$ (mean $\Delta G = -0.135$; Appendix A), providing evidence that D-study budgets have limited transferability to finer-grained scoring rubrics.

3.6 Model Facet Specification

Our target universe comprises open-weight, instruction-tuned models in the 7–9B parameter range (eight architecture families); results may not generalize to API-based, base, or $>9\text{B}$ models. Our eight purposively selected models are treated as a quasi-random facet (Meredith, 1993; Putnick and Bornstein, 2016). We report G under random-model assumptions, the more conservative specification: across five benchmarks, treating models as fixed yields G_{fixed} of 0.951–0.990 (mean 0.975) versus G_{random} of 0.863–0.896 (mean 0.882), a consistent increase of 0.077–0.113 (mean 0.093), because model $\times$ item interaction (30–36%) enters the relative error variance under random assumptions. Consequently, our D-study budgets are upper bounds—practitioners will not under-sample regardless of facet classification. The variance decomposition itself (Tables 4–5) is invariant to this choice. Leave-one-model-out analysis preserves component rankings across all five benchmarks, and Cohen’s κ pairwise agreement (mean $\kappa = 0.33$ –0.46) confirms genuine differential item functioning (Appendix A). Four-model subset consistency checks match the observed eight-model G within mean $|\Delta G| = 0.035$ across all five benchmarks (§4.5).

For reference, G_{model} (treating models as the object of measurement) under the recommended 3-prompt $\times$ 3-seed design with 200 items averages 0.79 across five benchmarks (range: 0.52 on ARC to 0.93 on HellaSwag). Model ranking reliability depends primarily on model $\times$ prompt interaction—ARC (8.0%) and GSM8K (8.5%) show lower G_{model} because prompt sensitivity reduces model discriminability. G_{item} remains the primary metric because item reliability is the direct optimization target for evaluation design; model reliability is a downstream consequence that improves with item count.

4 Experiments

All experiments use open-weight 7–9B parameter models evaluated with vLLM 0.9.2 on a single

NVIDIA RTX PRO 6000 GPU (96GB) with binary correctness scoring.

4.1 Single-Model Baseline

We first establish the variance structure within a single model (Llama-3.1-8B-Instruct) on MMLU to isolate methodological facet contributions before introducing cross-model variation.

Design. The G-study crosses six facets: 200 items $\times$ 3 numerical precisions (BF16, FP16, FP32) $\times$ 6 prompt templates $\times$ 6 random seeds $\times$ 3 temperature levels (0.0, 0.3, 0.7) $\times$ 4 answer orderings, yielding 259,200 binary observations. Variance components are estimated via Henderson Method I with bootstrap confidence intervals (1,000 resamples).

Results. Under binary scoring, item difficulty (*item_id*) accounts for 58.4% of total variance, with the residual term contributing 23.1%. All main facet effects are negligible ($<1\%$ each): precision 0.9%, prompt template, temperature, seed, and ordering each $<0.5\%$. The interesting structure appears in item-conditional interactions: *prompt* $\times$ *item* is the largest at 5.5%, followed by *seed* $\times$ *item* (4.7%), *precision* $\times$ *item* (3.3%), *temperature* $\times$ *item* (2.9%), and *ordering* $\times$ *item* (0.8%). This pattern indicates that methodological choices do not shift performance uniformly across items; instead, specific items are differentially sensitive to specific facets—prompt template variation produces the largest item-conditional effect.

The resulting generalizability coefficient is $G_{\text{item}} = 0.936$, indicating that 200 items under multi-facet conditions produce highly reliable rankings. A D-study projection shows $G \geq 0.80$ is achievable with as few as 55 items in this multi-facet design.

For text-exact-match scoring (requiring verbatim string matches), the variance structure shifts dramatically: temperature dominates at 79.7% of variance, and G_{item} collapses to 0.007. This metric-conditional divergence motivates our exclusive use of binary correctness for the remaining analyses.

4.2 Cross-Model Variance

Extending to eight models (BF16 subset, 460,800 records), Table 4 shows that item difficulty drops from 58.4% to 37.9%, while *model* $\times$ *item* interaction emerges at 30.2% (up to 36% across benchmarks, §4.3). The model main effect is 3–5% (MC)

Table 3: LOMO stability across five benchmarks (8 models $\times$ 5 benchmarks = 40 subsets).

Benchmark	G_{full}	CV (%)	Max $ \Delta G $	Most impactful
MMLU	0.896	0.49	0.019	InternLM2.5-7B
ARC	0.892	0.94	0.026	Llama-3.1-8B
HellaSwag	0.872	1.09	0.028	Llama-3.1-8B
GSM8K	0.863	1.25	0.030	Gemma-2-9B
MATH	0.888	1.20	0.029	InternLM2.5-7B

Table 4: Variance decomposition for the cross-model MMLU design (8 models, BF16, 460,800 records). Components estimated via Henderson Method I. $G_{\text{item}} = 0.896$.

Component	Var.%	95% CI
item_id	37.9	[33.1, 42.8]
<u>model</u> $\times$ <u>item</u>	<u>30.2</u>	[26.4, 34.1]
residual	25.0	[22.8, 27.3]
model	2.8	[1.2, 5.1]
model $\times$ prompt	1.6	[0.9, 2.5]
prompt $\times$ item	0.9	[0.5, 1.4]
temp. $\times$ item	0.6	[0.3, 1.0]
seed $\times$ item	0.5	[0.2, 0.8]
ordering $\times$ item	0.3	[0.1, 0.6]
all others (each)	<0.1	—

and 9–11% (FF)², confirming that models differ primarily in *which items* they find difficult rather than in overall ability level (contamination implications in §6). Leave-one-model-out analysis (Table 3) confirms this structure is stable: no single model removal changes G by more than 0.030 across all five benchmarks.

4.3 Exploratory: Format-Associated Variance Patterns

We extend to five benchmarks (three MC: MMLU, ARC-Challenge, HellaSwag; two FF: GSM8K, MATH) using exp-002 (8 models $\times$ BF16 $\times$ 2 temperatures $\times$ 6 prompts $\times$ 6 seeds $\times$ 4 orderings $\times$ 200 items). Table 5 reveals a format-associated pattern: MC benchmarks concentrate item-related variance (*item* + *model* $\times$ *item* DIF) at 68–72%, while FF benchmarks disperse it to 31–41% ($n = 5$; permutation $p = 0.10$; item-level $d = 1.73$, benchmark-level CI spans zero); we report this as a descriptive observation. This does not extend to item budgets: HellaSwag (MC, 118 items for $G \geq 0.80$) exceeds ARC (98) and MATH (101), showing format alone does not predict evaluation cost (Figure 5). The factorial design is not fully symmetric: ordering applies only to MMLU, and prompt templates serve format-specific func-

²Ranges rounded to nearest integer; exact values are 2.8–5.4% (MC) and 8.5–10.7% (FF) per Table 5.




Figure 3: Item difficulty heterogeneity ($\text{std}(p_i)$) versus item_id variance share across five benchmarks. MC benchmarks (filled) exhibit higher difficulty dispersion and item variance than FF benchmarks (open). Spearman $\rho = 0.82$.

Table 5: Format-associated variance structure across five benchmarks (exp-002, 8 models). Variance components (%) per benchmark; MC = multiple-choice, FF = free-form. Bold = dominant non-residual component. [†]For MC, reports prompt $\times$ item interaction (prompt main effect $< 1\%$); for FF, reports prompt main effect (item-conditional interactions absorbed into residual).

	MC			FF	
	MMLU	ARC	Hella.	GSM	MATH
item	37.9	37.9	34.3	16.9	24.0
mod. $\times$ item	30.2	33.7	36.0	14.6	17.3
model	2.8	3.1	5.4	10.7	8.5
prompt [†]	0.9	1.4	2.1	5.4	0.8
mod. $\times$ prmp	1.6	8.0	0.2	8.5	1.7
residual	25.0	15.4	21.4	39.9	43.9
all others	1.6	0.5	0.6	4.0	3.8

tions, potentially confounding format comparisons. Figure 3 shows that item difficulty heterogeneity ($\text{std}(p_i)$) predicts item variance share across benchmarks ($\rho = 0.82$), with MC benchmarks clustering at higher dispersion and item variance than FF benchmarks. Per-benchmark details—including HellaSwag and ARC exceptions, FF residual analysis (Table 12; Table 11)—are in Appendix D. All format-associated findings are conditional on binary scoring (§3.5).

4.4 Precision Separability

Under binary scoring, BF16 and FP16 are functionally interchangeable ($< 0.1\%$ interaction (Yuan et al., 2025)); under text-exact-match, a precision $\times$ temperature interaction emerges at 1.3%, consistent with §4.1. Precision need not be controlled for binary evaluation at BF16/FP16 (Shanmugavelu et al., 2024).




Figure 4: Item variance share (%) by model and benchmark. MC benchmarks (left) consistently show higher item variance (46–62%) than FF benchmarks (right, 18–32%) across all eight models, confirming the format-associated pattern is not driven by individual models.




Figure 5: Variance decomposition (Henderson Method I) across five benchmarks and eight models. Multiple-choice benchmarks share an item-dominated structure; HellaSwag is the exception where model $\times$ item (36.0%) exceeds item difficulty (34.3%). Free-form benchmarks (GSM8K, MATH) distribute variance across item, model, and interaction components with no single dominant source.

4.5 Design-Dependent Reproducibility Budgets

How many items does a given evaluation design need to achieve a target reliability? We translate variance components into item budgets via D-study projections.

Table 6 compares evaluation designs at the $G \geq 0.80$ threshold using single-model (Llama-3.1-8B-Instruct) MMLU variance components. The multi-facet design (1,296 conditions per item) achieves $G = 0.936$ at 200 items and requires only 55 items for $G \geq 0.80$. A single-condition design reaches only $G = 0.592$ at 200 items and requires 552 items—a 10.0 $\times$ budget ratio stable across subset sizes (CV = 1.9%).

Zhou et al. (2026)’s IRT-selected 116 items yield $G = 0.457$ under single-condition but $G = 0.895$ under multi-facet conditions, demonstrating that

Table 6: D-study reproducibility budgets under different evaluation designs (single-model, MMLU). Total evals = items $\times$ conditions.

Design	Cond.	$n_{G \geq .80}$	Total	G_{200}	G_{116}
Multi-facet	1,296	55	71,280	.936	—
Single-cond.	1	552	552	.592	.457
IRT+multi	1,296	—	150,336	—	.895 [‡]
IRT baseline	—	116	116	—	.457 [†]

[†] PSN-IRT 116 items (Zhou et al., 2026) under single-condition: $G = 0.457$.

[‡] Same 116 items under multi-facet: $G = 0.895$.




Figure 6: Effect of $|\Delta G| = 0.035$ estimation error on item budgets across five benchmarks. The $\sim 23\%$ budget increase is consistent across all benchmarks, indicating D-study budgets are robust to estimation uncertainty.

IRT item selection and G-theory protocol optimization are orthogonal and complementary. An internal consistency check using four-model subsets confirms projection stability ($|\Delta G| \leq 0.029$ for MC, ≤ 0.065 for FF; Table 9); the corresponding $\sim 23\%$ budget increase is consistent across all five benchmarks (Figure 6).

The design choice depends on the evaluation bottleneck: the full 1,296-condition design minimizes items (55, 71,280 evaluations/model); the recommended 3×3 design balances cost and reliability (178 items, 1,602 evaluations); single-condition requires 552 items. Cross-model projections show budget ratios of $11\text{--}22\times$ across benchmarks (bootstrap CIs all exceed $10\times$; Appendix B), with multi-facet budgets of $93\text{--}128$ items (Table 8).

Projection assumptions. Cross-model budgets increase to 93 items because model $\times$ item interaction enters relative error variance. G_{item} is stable across benchmarks (0.863–0.896, CV = 1.4%; LOMO in Appendix A). All projections assume the 2024–2025 model generation; the 72B extension (§4.6) provides preliminary support.

4.6 Scale Validation

To test whether the variance structure generalizes beyond 7–9B models, we conduct within-family and cross-family analyses. Qwen2.5-72B-Instruct on MMLU shows item difficulty at 83.2% ($G_{\text{item}} = 0.979$), with prompt $\times$ item (6.8%) as the

dominant interaction—the same ordering observed at 7B. A three-model crossed design (Qwen2.5 3B/7B/72B) preserves the model $\times$ item dominant interaction (38.5%) and item main effect (35.7%). Cross-family, Llama-3.1-70B-Instruct yields item difficulty at 74.5% ($G_{\text{item}} = 0.957$), confirming that item difficulty dominates across a $24\times$ parameter range and two architecture families.

Robustness checks. Three MMLU resamples yield CV = 3.1%; treating models as fixed increases G by 0.077–0.113, confirming conservative budgets (§3.6; seven additional analyses in Appendix A).

5 Conclusion

A recommended 3-prompt $\times$ 3-seed design reduces item requirements from 552 to 178 for $G \geq 0.80$, confirmed across five benchmarks under binary scoring. The variance structure is format-dependent: MC benchmarks concentrate 68–72% of variance in item-related components, while FF benchmarks disperse to 31–41%, suggesting that reproducibility budgets should be calibrated per benchmark format. All budgets assume binary scoring; extending to other model scales and continuous metrics is a priority. We recommend that benchmark designers adopt multi-facet validation (at minimum 3 prompts $\times$ 3 seeds) as standard practice, and that leaderboard operators average over prompt and seed conditions to improve ranking stability. Code, data, and a D-study calculator will be released upon publication.

6 Limitations

Our design covers eight 7–9B instruction-tuned models; preliminary 72B validation confirms consistent variance structure, but broader scale and temporal coverage is needed. The benchmark sample ($n = 5$) limits statistical power for format-associated observations ($p = 0.10$). All D-study budgets assume binary scoring; partial-credit analysis on GSM8K shows G drops from 0.86 to 0.77, suggesting limited transferability to finer-grained rubrics.

References

Xavier Bouthillier, Pierre Delaunay, Mirko Bronzi, Asya Trofimov, Brennan Nichyporuk, Justin Szeto, Nazanin Mohammadi Sepahvand, Edward Raff, Kshiti Madan, Vikram Voleti, Samira Ebrahimi Kahou,

- Vincent Michalski, Tal Arbel, Chris Pal, Gaël Varoquaux, and Pascal Vincent. 2021. [Accounting for variance in machine learning benchmarks](#). In *Proceedings of Machine Learning and Systems (MLSys)*.
- Robert L. Brennan. 2001. *Generalizability Theory*. Springer.
- Ngoc Thang Bui, Guergana Savova, and Lijun Wang. 2025. [Assessing the macro and micro effects of random seeds on fine-tuning large language models](#). In *Proceedings of IJCNLP-AACL*.
- Brian E. Clauser and Kathleen M. Mazor. 1998. [Using statistical procedures to identify differentially functioning test items](#). *Educational Measurement: Issues and Practice*, 17(1):31–44.
- Mostafa Dehghani, Yi Tay, Alexey A. Gritsenko, Zhe Zhao, Neil Houlsby, Fernando Diaz, Donald Metzler, and Oriol Vinyals. 2021. [The benchmark lottery](#). *arXiv preprint arXiv:2107.07002*.
- Jesse Dodge, Gabriel Ilharco, Roy Schwartz, Ali Farhadi, Hannaneh Hajishirzi, and Noah A. Smith. 2020. [Fine-tuning pretrained language models: Weight initializations, data orders, and early stopping](#). *arXiv preprint arXiv:2002.06305*.
- Bradley Efron and Robert J. Tibshirani. 1993. *An Introduction to the Bootstrap*. Chapman & Hall, New York.
- Leo Gao, Jonathan Tow, Baber Abbasi, Stella Biderman, Sid Black, Anthony DiPofi, Charles Foster, Laurence Golber, Jeffrey Hsu, Alain Le Noac’h, Haonan Li, Kyle McDonell, Niklas Muennighoff, Chris Ocqueateau, Jason Phang, Laria Reynolds, Hailey Schoelkopf, Aviya Skowron, Lintang Sutawika, and 5 others. 2024. [A framework for few-shot language model evaluation](#).
- Haiyang He and Thinking Machines Lab. 2025. [Defeating nondeterminism in LLM inference](#). Thinking Machines Lab: Connectionism.
- Peter Henderson, Riashat Islam, Philip Bachman, Joelle Pineau, Doina Precup, and David Meger. 2018. [Deep reinforcement learning that matters](#). In *Proceedings of the AAAI Conference on Artificial Intelligence*.
- Paul W. Holland and Howard Wainer, editors. 1993. *Differential Item Functioning*. Lawrence Erlbaum Associates, Hillsdale, NJ.
- Angie Hua, Kevin Tang, Chaoqun Gu, Jing Gu, Edmund Wong, and Yilin Qin. 2025. [Flaw or artifact? rethinking prompt sensitivity in evaluating LLMs](#). In *Proceedings of the Conference on Empirical Methods in Natural Language Processing (EMNLP)*.
- Zhehan Jiang and William Skorupski. 2018. [A Bayesian approach to estimating variance components within a multivariate generalizability theory framework](#). *Behavior Research Methods*, 50.
- Ryan Othniel Kearns. 2026. [Quantifying construct validity in large language model evaluations](#). *arXiv preprint arXiv:2602.15532*.
- Alex Kipnis, Konstantinos Voudouris, Luca Schulze Buschoff, and Eric Schulz. 2025. [metabench – a sparse benchmark of reasoning and knowledge in large language models](#). In *Proceedings of the International Conference on Learning Representations*.
- John P. Lalor, Hao Wu, and Hong Yu. 2019. [Learning latent parameters without human response patterns: Item response theory with artificial crowds](#). In *Proceedings of the Conference on Empirical Methods in Natural Language Processing and International Joint Conference on Natural Language Processing (EMNLP-IJCNLP)*.
- Guangming Li. 2023. [Which method is optimal for estimating variance components and their variability in generalizability theory? Evidence from a set of unified rules for bootstrap method](#). *PLoS ONE*.
- Percy Liang, Rishi Bommasani, Tony Lee, Dimitris Tsipras, Dilara Soylu, Michihiro Yasunaga, Yian Zhang, Deepak Narayanan, Yuhuai Wu, Ananya Kumar, and 1 others. 2023. [Holistic evaluation of language models](#). *arXiv preprint arXiv:2211.09110*.
- Lovish Madaan, Samarth Bhatia, Marco Tulio Ribeiro, and Dan Roth. 2024. [Quantifying variance in evaluation benchmarks](#). *arXiv preprint arXiv:2406.10229*.
- Felipe Maia Polo, Lucas Weber, Leshem Choshen, Yuekai Sun, Gongjun Xu, and Mikhail Yurochkin. 2024. [tinyBenchmarks: Evaluating LLMs with fewer examples](#). In *Proceedings of the International Conference on Machine Learning (ICML)*.
- William Meredith. 1993. [Measurement invariance, factor analysis and factorial invariance](#). *Psychometrika*, 58(4):525–543.
- Alberto Messina and Sara Scotta. 2026. [Introducing background temperature to characterise hidden randomness in large language models](#). *arXiv preprint arXiv:2604.22411*.
- Solomon Messing. 2026. [Hidden measurement error in LLM pipelines distorts annotation, evaluation, and benchmarking](#). *arXiv preprint arXiv:2604.11581*.
- Cilia Mushquash and Brian P. O’Connor. 2006. [SPSS and SAS programs for generalizability theory analyses](#). *Behavior Research Methods*.
- Yotam Perlitz, Elron Bandel, Ariel Gera, Ofir Arviv, Liat Ein-Dor, Eyal Shnarch, Noam Slonim, Michal Shmueli-Scheuer, and Leshem Choshen. 2023. [Efficient benchmarking of language models](#). *arXiv preprint arXiv:2308.11696*.
- Pouya Pezeshkpour and Estevam Hruschka. 2024. [Large language models sensitivity to the order of options in multiple-choice questions](#). In *Findings of*

the North American Chapter of the Association for Computational Linguistics (NAACL).

Diane L. Putnick and Marc H. Bornstein. 2016. [Measurement invariance conventions and reporting: The state of the art and future directions for psychological research](#). *Developmental Review*, 41:71–90.

Pedro Rodriguez, Joe Barrow, Alexander Hoyle, John P. Lalor, Robin Jia, and Jordan Boyd-Graber. 2021. [Evaluation examples are not equally informative: How should that change NLP leaderboards?](#) In *Proceedings of the Association for Computational Linguistics (ACL)*.

Angelika Romanou, Marah Ibrahim, Candace Ross, Chantal Shaib, Kerem Oktar, Samuel J. Bell, Anaelia Ovalle, Jesse Dodge, Antoine Bosselut, Koustuv Sinha, and Adina Williams. 2026. [Brittlebench: Quantifying LLM robustness via prompt sensitivity](#). *arXiv preprint arXiv:2603.13285*.

Melanie Sclar, Yejin Choi, Yulia Tsvetkov, and Alane Suhr. 2024. [Quantifying language models’ sensitivity to spurious features in prompt design or: How I learned to start worrying about prompt formatting](#). In *Proceedings of the International Conference on Learning Representations (ICLR)*.

Sanjif Shanmugavelu, Mathieu Taillefumier, Christopher Culver, Oscar Hernandez, Mark Coletti, and Ada Sedova. 2024. [Impacts of floating-point non-associativity on reproducibility for HPC and deep learning applications](#). In *SC’24 Workshops (IEEE)*.

Richard J. Shavelson and Noreen M. Webb. 1991. *Generalizability Theory: A Primer*. SAGE Publications.

Richard J. Shavelson, Noreen M. Webb, and Glenn L. Rowley. 1989. [Generalizability theory](#). *American Psychologist*.

Francis Tuerlinckx, Frank Rijmen, Geert Verbeke, and Paul De Boeck. 2006. [Statistical inference in generalized linear mixed models: A review](#). *British Journal of Mathematical and Statistical Psychology*.

Haoran Ye, Jing Jin, Yuhang Xie, Xin Zhang, and Guojie Song. 2025. [Large language model psychometrics: A systematic review of evaluation, validation, and enhancement](#). *arXiv preprint arXiv:2505.08245*.

Juntao Yuan, Haoyang Li, Xingcheng Ding, Weijie Xie, Yue-Jiao Li, Wentao Zhao, Kaiyan Wan, Jiayi Shi, Xiaolu Hu, and Zhouhan Liu. 2025. [Understanding and mitigating numerical sources of nondeterminism in LLM inference](#). In *Advances in Neural Information Processing Systems (NeurIPS)*. Oral presentation.

Hao Zhou, Huifeng Huang, Zhe Zhao, Long Han, Haobo Wang, Ke Chen, Min Yang, Wanli Bao, Jiaxin Dong, Bo Xu, Conghui Zhu, Hao Cao, and Tiejun Zhao. 2026. [Lost in benchmarks? rethinking large language model benchmarking with item response theory](#). In *Proceedings of the AAAI Conference on Artificial Intelligence*. Oral presentation.

A Additional Robustness Analyses

We report nine robustness analyses that verify the stability of the variance decomposition under alternative analytical choices.

GLMM complement. A generalized linear mixed model (GLMM) with marginal logit link produces variance component rankings that perfectly preserve all seven grouped component orderings ($\rho = 1.0$) relative to Henderson Method I (Jiang and Skorupski, 2018; Tuerlinckx et al., 2006). This confirms that the linear variance decomposition is robust to the binary outcome distribution.

Accuracy-range stability. The `item_id` variance share remains stable across subsets with mean accuracy from 0.62 to 0.72 ($r = -0.048$). This rules out distortion from binary floor or ceiling effects at observed accuracy levels.

Estimator robustness. We compare Henderson Method I with restricted maximum likelihood (REML-EM, 500 iterations, practical convergence $\Delta < 0.003$ percentage points). All variance components differ by < 0.5 percentage points between estimators. The largest discrepancy occurs on the residual component (0.47pp); all substantive components (`item`, `model`×`item`, `model`) differ by < 0.3 pp. Henderson Method I is preferred for our balanced designs because it produces unbiased estimates without distributional assumptions (Brennan, 2001, Ch. 12).

Probit simulation. A probit simulation matched to our 1,296-condition design confirms that Henderson Method I recovers variance component rankings from dichotomized data (Spearman $\rho = 0.732$ between binary and continuous rankings, excluding residual; `item_id` dominant in all 100 replications; top-5 ordering correct in 79% of replications). Bias analysis by difficulty stratum shows $|\text{bias}| = 0.28\text{--}0.37$ for items with $p \in [0.5, 0.8]$, where the Bernoulli variance ceiling $p(1-p)$ is largest. This analysis covers the observed accuracy range ($p \in [0.3, 0.9]$); extreme-difficulty items ($p < 0.3$ or $p > 0.9$) are rare in our sample ($< 8\%$ of items) and may exhibit larger deviations. Henderson Method I absorbs this Bernoulli-inherent variance into the residual, explaining inflated bias near medium difficulty; systematic component ordering is robustly recovered across all replications. Binary scoring inflates residual variance by absorbing the Bernoulli $p(1-p)$ ceiling, which enters σ_{rel}^2

and increases D-study item requirements. Within a matched probit simulation, binarization increases D-study budgets by approximately 36% (99 → 135 items for $G \geq 0.80$), driven by residual inflation (16% → 41%) and item variance compression (50% → 32%). Facet $\times$ item interactions—the components entering σ_{rel}^2 —are best preserved under binarization (MAPE = 12.8%, range 8.7–17.5%), while residual variance inflates as expected from Bernoulli $p(1-p)$ noise (MAPE = 156%). D-study budgets derived from binary data are therefore conservative: continuous-score designs would require equal or fewer items for the same G threshold. Null calibration (ratio = 1.0008) confirms that this conservatism does not distort relative component magnitudes.

CI method comparison. Bootstrap (1,000 item-level resamples with replacement) and delete-1 jackknife confidence intervals agree closely: width ratios range from 0.96 to 1.04 across all variance components. Bootstrap CIs capture item-sampling variability conditional on these 8 specific models; model-level stability is assessed separately via LOMO (below). The agreement indicates that the dependence structure introduced by item-level bootstrap resampling (which creates ~ 127 unique items per 200-item resample) does not materially distort interval estimates. Stratified resampling by subject (10 subjects $\times$ 20 items) produces CIs within 4% of simple item-level bootstrap, confirming that the within-subject dependence structure does not materially affect interval estimates.

Facet specification sensitivity. Treating temperature and ordering as fixed (rather than random) facets changes G_{item} by only 0.37 percentage points (all-random: 0.820 vs. fixed temperature+ordering: 0.824). The small difference reflects the negligible temperature $\times$ item (0.65%) and ordering $\times$ item (0.54%) variance shares. The random-facet assumption is therefore robust to facet classification.

Leave-one-model-out (LOMO). Across all 40 leave-one-model-out subsets (8 models $\times$ 5 benchmarks), G_{item} ranges from 0.833 to 0.919, with CV $\leq 1.25\%$ per benchmark. No single model removal changes G by more than 0.030 (GSM8K, removing Gemma-2-9B). Item-related variance rankings are preserved in all 40 subsets. Table 3 (main text) reports per-benchmark stability. This supports treating models as exchangeable within the target universe of 7–9B instruction-tuned models (§3.6).

Inter-model agreement. We computed pairwise Cohen’s κ and Fleiss’ κ across all model pairs on each benchmark. Mean pairwise κ ranges from 0.326 (MATH) to 0.461 (MMLU); Fleiss’ κ ranges from 0.309 (GSM8K) to 0.456 (MMLU). HelLaSwag uses six models ($\kappa = 0.332$) due to missing OLMo and Yi data. The fair-to-moderate agreement independently confirms that the model $\times$ item interaction (30–36% of variance) reflects genuine differential item functioning: models disagree substantially on which items are difficult, consistent with prior DIF findings in educational measurement (Holland and Wainer, 1993; Clauser and Mazor, 1998).

Scoring metric sensitivity. To probe metric conditionality, we compared binary correctness with step-level partial credit on GSM8K using the full 8-model crossed design (200 items, 2,304 conditions per item). Under partial credit, G_{item} drops from 0.863 to 0.768, with residual variance rising from 39.9% to 64.9% and item variance falling from 16.9% to 10.3%. All eight models individually show $G_{\text{partial}} < G_{\text{binary}}$ (mean $\Delta G = -0.135$, range -0.025 to -0.309), with ΔG magnitude moderated by model accuracy: higher-accuracy models have fewer incorrect responses whose scores shift under partial credit (e.g., OLMo at 70.1% accuracy shows $\Delta G = -0.025$, while Llama at 49.0% shows $\Delta G = -0.309$). The mechanism is clear: 48.1% of incorrect responses have numerically correct arithmetic steps but incorrect problem setup, receiving partial score = 1.0 despite binary score = 0—this compresses item discrimination and inflates residual variance. Binary scoring amplifies prompt sensitivity by collapsing reasoning quality into pass/fail, confirming that all findings in this paper are metric-conditional.

Facet specification sensitivity: random vs. fixed model. Table 7 reports per-benchmark generalizability coefficients under random and fixed model specifications. The consistent ΔG of 0.077–0.113 confirms that random-model budgets are strictly more conservative across all benchmarks.

Null calibration. We generate 100 Bernoulli-distributed datasets with item difficulties matched to the empirical distribution but no systematic facet effects. The observed-to-null variance ratio is 1.0008, confirming that Henderson Method I does not introduce systematic bias when applied to binary outcomes. The null simulation also verifies

Table 7: Generalizability coefficients under random vs. fixed model facet specification across five benchmarks.

Benchmark	G_{random}	G_{fixed}	ΔG
MMLU	0.896	0.984	+0.088
ARC	0.892	0.990	+0.098
HellaSwag	0.872	0.984	+0.113
GSM8K	0.863	0.951	+0.089
MATH	0.888	0.965	+0.077
Mean	0.882	0.975	+0.093

that the item variance share (58.4%) is not an artifact of the Bernoulli $p(1-p)$ ceiling: the null produces item variance of 0.08%, three orders of magnitude below the observed value. Across all 100 replications, the same three components (item, model $\times$ item, residual) emerge as the top three (set consistency 100%), with Kendall’s $\tau = 0.81$ between binary and continuous rankings.

Reduced crossing validation. To verify that the reduced crossing in exp-002 (BF16 only $\times$ 2 temperatures) does not introduce confounding, we extracted the corresponding subset from the fully crossed exp-001 design. The generalizability coefficient differs by only 0.0008 ($G = 0.9361$ vs. 0.9369), and all non-precision variance components differ by less than 2 percentage points. The 10pp increase in item variance share (58.4% $\rightarrow$ 68.5%) is fully explained by the removal of precision-related components ($\sim 4.5\%$ total), which redistributes across remaining sources. This confirms that the exp-002 design preserves the variance structure of the full crossing.

B Cross-Model D-Study Projections

Table 8 extends the single-model D-study (Table 6 in main text) to per-benchmark cross-model projections with bootstrap confidence intervals (1,000 item-level resamples, 8 models, 2,304 conditions per item). The efficiency ratio is format-associated: MC benchmarks show 11–13 $\times$ (mean 12.0 $\times$ [11.6, 12.5]), while FF benchmarks show 21–22 $\times$ (mean 21.5 $\times$ [20.4, 22.8]).

The higher item requirements relative to the single-model analysis (55 items in Table 6) reflect the additional model $\times$ item interaction variance that enters σ_{rel}^2 when models are crossed as a facet. The format-associated ratio pattern—FF benchmarks showing $\sim 2\times$ the MC ratio—is consistent with the distributed variance structure observed in §4.3, where FF benchmarks allocate more variance to model and model $\times$ prompt components.

Table 8: Cross-model D-study reproducibility budgets with 95% bootstrap CIs. $n_{G \geq 0.80}$: items needed. Multi = multi-facet (2,304 cond.); Total = items $\times$ 2,304. Single = single-condition (total evals = item count).

	Multi-facet (2,304 cond.)			Single-cond.		Ratio [95% CI]
	Items	95% CI	Total	Items	95% CI	
<i>Multiple-choice</i>						
MMLU	93	[76, 116]	214K	1,216	[1004, 1491]	13.1 [12.5, 13.7]
ARC	98	[71, 139]	226K	1,078	[793, 1512]	11.0 [10.7, 11.4]
Hella.	118	[93, 151]	272K	1,403	[1120, 1790]	11.9 [11.5, 12.3]
<i>Free-form</i>						
GSM8K	128	[103, 162]	295K	2,768	[2230, 3531]	21.6 [20.7, 23.0]
MATH	101	[86, 122]	233K	2,160	[1853, 2631]	21.4 [20.2, 22.5]

Internal consistency check. Table 9 reports consistency checks from four exp-001 models (DeepSeek, Llama, Mistral, Qwen3) to the full eight-model design across all five benchmarks. MC benchmarks show $|\Delta G| \leq 0.029$ ($< 3.5\%$ error), while FF benchmarks show larger deviations ($|\Delta G| \leq 0.065$, $\sim 7\%$) due to limited model diversity in the four-model subset. FF predictions consistently underestimate actual G (conservative); MC predictions show mixed directions with small deviations ($|\Delta G| \leq 0.029$).

Table 9: Internal consistency: four-model subset vs. eight-model G_{item} . FF subsets underestimate actual G (conservative); MC shows mixed directions with small deviations.

Benchmark	Predicted G	Actual G	$ \Delta G $	Error%
<i>Multiple-choice</i>				
MMLU	0.888	0.896	0.009	0.96
ARC	0.904	0.892	0.012	1.38
HellaSwag	0.901	0.872	0.029	3.33
<i>Free-form</i>				
GSM8K	0.803	0.863	0.059	6.88
MATH	0.823	0.888	0.065	7.30
Mean	—	—	0.035	3.97

Budget sensitivity to estimation error. The effect of $|\Delta G| = 0.035$ estimation error on item budgets is shown in Figure 6 (main text). The percentage increase is remarkably stable (23.0–24.5%, mean 23.2%), because the relative budget adjustment $\Delta n/n$ depends on the G -formula ratio rather than on absolute variance components. Absolute Δn ranges from 13 (ARC) to 67 (GSM8K), scaling with the baseline budget.

Cross-benchmark G stability. G_{item} values cluster in a narrow band across all five benchmarks (0.863–0.896, CV = 1.4%), despite qualitatively different variance structures: MC benchmarks concentrate 68–72% in item-related components while

GSM8K disperses variance across item (17%), model (11%), and interaction (15%) sources (Table 5). This stability implies that D-study budget recommendations calibrated on one benchmark provide reasonable order-of-magnitude guidance for others. However, the 0.033 range in G_{item} translates to a 71–128 item range for multi-facet $G \geq 0.80$ (Table 8), and the variance compositions differ substantially—benchmark-specific G-studies remain necessary to identify dominant variance sources and optimize condition selection.

C Intermediate D-Study Designs

Table 10 shows how reliability improves as evaluation designs increase from single-condition to full multi-facet crossing. The steepest gain occurs at the “minimal upgrade” level (2 prompts): G_{200} jumps from 0.592 to 0.692 at $2\times$ cost. A minimal sufficient design for $G \geq 0.80$ is 3 prompts $\times$ 3 seeds (9 conditions, $G_{200} = 0.818$, $n_{G \geq 0.80} = 178$). Beyond the “standard” level (32 conditions), adding orderings provides negligible benefit (ordering $\times$ item $< 1\%$), confirming that prompt and seed diversity are the primary reliability drivers.

Table 10: Intermediate D-study designs (single-model, MMLU). Designs are ordered by increasing conditions per item; G_{200} and $n_{G \geq 0.80}$ are D-study projections from exp-001 variance components.

Design	Conditions	G_{200}	$n_{G \geq .80}$
Single-condition	1	.592	552
Minimal upgrade (2p)	2	.692	357
Easy upgrade (2p $\times$ 2s)	4	.766	245
Moderate (3p $\times$ 3s)	9	.818	178
Standard (4p $\times$ 4s $\times$ 2t)	32	.869	121
Enhanced (+4 orderings)	72	.871	119
Full multi-facet	1,296	.936	55

D Format-Associated Variance: Per-Benchmark Details

This section supplements the compressed format-associated analysis in §4.3 with per-benchmark details.

MC benchmark exceptions. HellaSwag is the sole MC case where model $\times$ item slightly exceeds item (36.0% vs. 34.3%), though their sum (70.3%) remains in the MC range (68–72%). Among MC benchmarks, ARC-Challenge shows an anomalously high model $\times$ prompt interaction (8.0% vs. $< 2\%$ for MMLU and HellaSwag), suggesting that

science-domain items are differentially sensitive to option formatting across model families.

FF benchmark patterns. Free-form benchmarks differ markedly from MC. On GSM8K, variance distributes across item (16.9%), model $\times$ item (14.6%), model (10.7%), and model $\times$ prompt (8.5%); prompt main effect is 5.4%. MATH shows an intermediate pattern: item leads at 24.0% with model $\times$ item at 17.3%, but residual variance is substantially larger (43.9%). Item-level regression reveals that FF residual variance is predominantly explained by the Bernoulli variance ceiling $p(1-p)$ ($R^2 = 0.47\text{--}0.86$; Appendix G): medium-difficulty items maximize binary outcome variance regardless of facet structure, so the elevated FF residuals reflect measurement resolution rather than unmodeled substantive sources.

Difficulty dispersion as variance predictor. Difficulty dispersion ($\text{std}(p_i)$) strongly predicts item variance share ($\rho = 0.82$, $n = 5$ benchmarks; Appendix E; reasoning-heavy subjects show G_{item} as low as 0.614, Appendix F). Format constrains the noise floor, but difficulty heterogeneity determines variance magnitude within it.

Item-level bootstrap stability. Bootstrap resampling (1,000 item resamples per benchmark) confirms that the MC–FF separation is robust to item selection: 95% CIs for item-related variance (item + model $\times$ item) are [65.2, 75.0]% (MC) and [28.0, 44.2]% (FF), with no overlap (minimum gap 21.0 percentage points).

E Cross-Benchmark Difficulty Heterogeneity

Item difficulty heterogeneity is shown in Figure 3 (main text). MC benchmarks cluster at higher dispersion (0.27–0.31) and higher item variance (34–38%), while FF benchmarks occupy the lower-left region (0.22–0.23, 17–24%). The cross-benchmark Spearman $\rho = 0.82$ ($p = 0.089$, $n = 5$) is consistent in direction with the within-MMLU correlation ($\rho = 0.806$, $p = 0.005$, $n = 10$ subjects); these are independent analyses at different granularities (benchmarks vs. subjects within MMLU), providing multi-scale support for the item difficulty heterogeneity mechanism discussed in §4.3.

F MMLU Subject-Level Variance Decomposition

Table 11 reports per-subject variance decomposition for 10 MMLU subjects (20 items each, 8 models, 2,304 conditions per item). G_{item} ranges from 0.614 (*high_school_mathematics*) to 0.930 (*international_law*), with mean 0.843. Knowledge-intensive subjects (*international_law*, *clinical_knowledge*, *world_religions*) show $\text{item_id} > 40\%$, while reasoning-heavy subjects (*high_school_mathematics*, *abstract_algebra*) show $\text{model} \times \text{item} > \text{item_id}$, consistent with models differing more in reasoning strategies than in factual recall. The *high_school_mathematics* item_id of 10.4% referenced in §4.3 is the lowest across all subjects, below even GSM8K (16.9%).

Table 11: Subject-level variance decomposition within MMLU (20 items per subject, 8 models, 2,304 conditions per item). Sorted by G_{item} descending.

Subject	Acc.	item%	mod. $\times$ it.%	G_{item}
<i>international_law</i>	.717	44.4	24.1	.930
<i>clinical_knowledge</i>	.581	47.5	30.0	.923
<i>world_religions</i>	.826	42.2	27.1	.910
<i>computer_security</i>	.740	37.0	24.9	.908
<i>college_physics</i>	.436	31.6	30.1	.876
<i>philosophy</i>	.782	33.0	36.1	.868
<i>abstract_algebra</i>	.457	24.1	32.8	.822
<i>us_foreign_policy</i>	.864	20.4	36.9	.802
<i>professional_medicine</i>	.671	19.9	41.7	.778
<i>high_school_math</i>	.385	10.4	41.9	.614
Mean	—	31.1	32.6	.843

G FF Residual Variance Analysis

To investigate whether the elevated residual variance in free-form benchmarks (GSM8K 39.9%, MATH 43.9%) reflects unmodeled substantive sources or measurement properties of binary scoring, we regress item-level residual variance on item characteristics across all five benchmarks (Table 12).

The dominant predictor is the quadratic difficulty effect ($p < 10^{-5}$ on all benchmarks): items near $p = 0.5$ show maximal residual variance, while items near the floor or ceiling show minimal residual. This inverted-U pattern directly mirrors the Bernoulli variance ceiling $p(1-p)$. Item content features (number of reasoning steps, answer length) are not significant predictors for GSM8K ($p > 0.38$), confirming that the difficulty-residual relationship is statistical rather than content-driven.

FF benchmarks show substantially higher R^2 (mean 0.66) than MC benchmarks (mean 0.25), consistent with FF items spanning a wider difficulty range. The MC-FF contrast in residual magnitude is therefore a mathematical consequence of the difficulty distribution under binary scoring, not evidence of additional unmodeled variance sources in free-form evaluation.

Table 12: Item-level residual variance regression. Difficulty² = quadratic (inverted-U) term reflecting the Bernoulli $p(1-p)$ ceiling. GSM8K additionally tests n_{steps} and answer_length (both $p > 0.38$).

Benchmark	R^2 (full)	R^2 (diff. only)	Diff. ² p
<i>Multiple-choice</i>			
MMLU	0.377	0.347	<0.001
ARC	0.221	0.204	<0.001
HellaSwag	0.156	0.140	<0.001
<i>Free-form</i>			
GSM8K	0.468	0.449	<0.001
MATH	0.856	0.845	<0.001

H Model and Prompt Details

Models. Table 13 lists the eight models with their HuggingFace identifiers.

Table 13: Models used in all experiments.

Short name	HuggingFace identifier	Params
Llama-3.1-8B	meta-llama/Llama-3.1-8B-Instruct	8B
Gemma-2-9B	google/gemma-2-9b-it	9B
Mistral-7B-v0.3	mistralai/Mistral-7B-Instruct-v0.3	7B
Qwen3-8B	Qwen/Qwen3-8B	8B
Yi-1.5-9B	01-ai/Yi-1.5-9B-Chat	9B
OLMo-2-7B	allenai/OLMo-2-1124-7B-Instruct	7B
InternLM2.5-7B	internlm/internlm2_5-7b-chat	7B
DeepSeek-7B	deepseek-ai/deepseek-llm-7b-chat	7B

Prompt templates. Our six prompt templates follow a surface perturbation design inspired by BrittleBench (Romanou et al., 2026). For multiple-choice benchmarks, templates vary along three dimensions: option layout (lettered vs. numbered), instruction wording (direct vs. contextual), and answer prompt cue. For free-form benchmarks, templates vary chain-of-thought elicitation strategy. All six templates are semantically equivalent—they request the same task with different surface forms—to isolate formatting effects from content effects. ARC-Challenge and HellaSwag use the same six surface patterns as MMLU with domain-appropriate system instructions; MATH mirrors the GSM8K pattern with `\boxed{ }` answer formatting.

MMLU templates (representative MC format):

1112	1. The following are	5. Work through this math	1159
1113	multiple choice questions	question carefully.\nQ:	1160
1114	(with answers) about	{question}\nWork:	1161
1115	{subject}.\n\n{question}\nA.	6. Math problem:\n{question}\nLet's	1162
1116	... D. ...\nAnswer:	solve this step by step.	1163
1117	2. Answer the following {subject}		
1118	questions by selecting the	I Use of AI Assistants	1164
1119	correct option.\n\nQuestion:	We used Claude (Anthropic) as an AI coding and	1165
1120	{question}\n(A) ... (B) ...	writing assistant throughout this work. Specifi-	1166
1121	(C) ... (D) ...\nCorrect	cally:	1167
1122	answer:		
1123	3. Below are {subject}	• Code: drafting experiment scripts, data-	1168
1124	multiple-choice problems.	processing pipelines, and analysis utilities; de-	1169
1125	Pick the right letter.\nQ:	bugging and refactoring.	1170
1126	{question}\n A) ... B) ...	• Writing: drafting and revising paper sections,	1171
1127	C) ... D) ...\nA:	polishing prose, and reformatting L ^A T _E X.	1172
1128	4. Answer the following {subject}	• Literature: summarizing related papers to	1173
1129	question by selecting A, B,	inform related-work section.	1174
1130	C, or D.\n{question}\n- A.		
1131	... - B. ... - C. ... - D.		
1132	...\nAnswer:		
1133	5. Below is a {subject} exam		
1134	question with four possible		
1135	answers. Choose the correct		
1136	one.\n{question}\nOption A:		
1137	... Option B: ... Option C:		
1138	... Option D: ...\nCorrect:		
1139	6. Question ({subject}):\n{question}\nChoices:		
1140	(A) ... (B) ... (C) ... (D)		
1141	...\nYour answer:		
1142	<i>GSM8K templates (representative FF format):</i>		
1143	1. Solve the following math		
1144	problem step by step.\nQ:		
1145	{question}\nA: Let's think		
1146	step by step.		
1147	2. Answer the following		
1148	math word problem. Show		
1149	your work.\n\nQuestion:		
1150	{question}\n\nSolution:		
1151	3. Solve each problem below.		
1152	Write your reasoning, then		
1153	give the final answer		
1154	after ####.\n\nProblem:		
1155	{question}\n\nReasoning:		
1156	4. You are a math tutor. Solve		
1157	the problem and show your		
1158	steps.\n{question}\n\nSteps:		